



Journal of Statistical Software

MMMMMM YYYY, Volume VV, Issue II.

doi: 10.18637/jss.v000.i00

fastsae: High-Performance Frequentist and Bayesian Small Area Estimation in R

Ridson Al Farizal P. 
Politeknik Statistika STIS

Azka Ubaidillah 
Politeknik Statistika STIS

Abstract

Small area estimation (SAE) plays a pivotal operational role in official statistics, enabling the production of reliable disaggregated socioeconomic, health, and demographic indicators when domain-specific sample sizes are small or zero. However, practitioners in national statistical offices (NSOs) face a triple computational and operational barrier: frequentist spatial and spatio-temporal models with bootstrap mean squared error (MSE) estimation suffer from severe memory bottlenecks and slow runtimes in interpreted languages; non-Gaussian indicators (such as rates, skewed expenditures, and counts) traditionally require computationally demanding Markov chain Monte Carlo (MCMC) algorithms that impede routine production; and model estimates often fail to aggregate to published higher-level totals, violating institutional consistency. We introduce **fastsae**, an R package engineered for high-performance frequentist and Bayesian small area estimation. On the frequentist side, **fastsae** implements optimized C++17 routines via **RcppArmadillo** paired with OpenMP multi-threading, accelerating the classical Fay-Herriot, spatial Fay-Herriot, spatio-temporal Fay-Herriot, and unit-level Battese-Harter-Fuller models alongside their analytical and bootstrap MSE estimators. On the Bayesian side, **fastsae** provides a unified empirical best prediction (EBP) interface powered by Integrated Nested Laplace Approximations (**INLA**) across six distribution families (Gaussian, Beta, Gamma, Binomial, Poisson, Negative Binomial). Crucially, **fastsae** introduces exact survey dispersion mappings that respect design-based direct sampling variances and accommodates modern spatial priors (BYM, scaled BYM2, Leroux CAR, and SLM) as well as Knorr-Held space-time interaction dynamics. To bridge estimation and dissemination, the package features a two-stage hierarchical benchmarking engine for simultaneous provincial and national calibration, an automated reliability classification engine aligned with NSO standards, multi-model concordances, automated **ggplot2** spatial mapping, and multi-sheet spreadsheet exports. Extensive benchmarks and finite population Monte Carlo simulations demonstrate multi-fold computational speedups (up to 364× for Fay-Herriot, 80× for Spatial Fay-Herriot, and 544× for spatio-temporal Beta) and peak RAM reductions up to 99.9% without loss of inferential precision.

Keywords: small area estimation, Fay-Herriot, empirical best prediction, INLA, OpenMP, Rcpp, hierarchical benchmarking, spatial statistics, R.

1. Introduction

In modern evidence-based governance, policy interventions are increasingly decentralized. Monitoring the United Nations Sustainable Development Goals (SDGs), allocating municipal fiscal transfers, targeting social safety net subsidies, and responding to localized disease outbreaks all require dependable statistical indicators at highly disaggregated geographic administrative levels (Rao and Molina 2015; Pfeffermann 2013). However, national statistical offices (NSOs) face structural sample size limitations. Large-scale household surveys—such as the American Community Survey (ACS) in the United States, the European Union Statistics on Income and Living Conditions (EU-SILC), and the National Socioeconomic Survey (Susenas) in Indonesia—are powered to yield acceptable sampling errors only at national or broad regional levels. When partitioned into sub-domains (e.g., municipalities, districts, or health zones), domain-specific sample sizes become sparse or zero, causing direct survey estimators to suffer from high sampling variance and unacceptable coefficients of variation (Burgard, Krause, and Merkle 2020; Pratesi 2016).

Small area estimation (SAE) overcomes these limitations by “borrowing strength” across domains and time periods via linking models that combine survey responses with auxiliary administrative registers, census aggregates, or geospatial layers (Rao and Molina 2015). The cornerstone of this literature remains the area-level Fay-Herriot (FH) model (Fay and Herriot 1979), extended to account for spatial dependencies via Simultaneous Autoregressive (SAR) structures (Petrucci and Salvati 2006; Pratesi and Salvati 2008) and temporal dynamics via panel autoregressive processes (Marhuenda, Molina, and Morales 2013). In the presence of micro-level auxiliary data, the nested error regression model of Battese, Harter, and Fuller (1988) (BHF) provides unit-level empirical best linear unbiased predictors (EBLUP).

Despite these theoretical advances, operationalizing small area models within production environments in NSOs has been hindered by three structural bottlenecks:

1. **Frequentist Computational and Memory Bottlenecks:** Existing implementations of spatial and spatio-temporal Fay-Herriot models rely on interpreted languages. Estimating model variance components via Restricted Maximum Likelihood (REML) requires repeated inversions of $D \times D$ or $DT \times DT$ covariance matrices. When evaluating Mean Squared Error (MSE) through parametric or non-parametric bootstrap resampling (González-Manteiga, Lombardía, Molina, Morales, and Santamaría 2008), runtimes escalate to hours. Crucially, memory allocation for dense covariance arrays grows quadratically ($O(D^2)$ to $O(D^2T^2)$), triggering out-of-memory (OOM) crashes on large national grids ($D \geq 500$).
2. **Distributional Inflexibility and MCMC Friction in Bayesian SAE:** Real-world welfare indicators are rarely Gaussian. Rates and proportions (poverty headcounts, unemployment, stunting) are strictly bounded in $(0, 1)$, while expenditures exhibit severe positive skewness (Gamma), and demographic events represent discrete counts (Poisson, Negative Binomial) (Janicki 2020; Fabrizi, Ferrante, and Trivisano 2011; Boubeta, Lombardía, and Morales 2016). Under non-Gaussian likelihoods, linear EBLUP is invalid, requiring Empirical Best Prediction (EBP) within a Hierarchical Bayesian (HB) framework. Currently, Bayesian SAE software relies on Markov Chain Monte Carlo (MCMC) sampling via **Stan** or **JAGS** (De Nicolò and Gardini 2022). MCMC sampling requires lengthy runtimes, complex convergence diagnostics ($\hat{R}$, trace plots), and often

Feature / Package	fastsae	sae	emdi	tipsae	mcmc-sae	SUMMER
Computational Engine	C++17 / INLA	Pure R	R / nlme	Stan HMC	C++ / MCMC	INLA
Frequentist EBLUP (FH/SFH)	✓	✓	✓	×	×	×
Spatio-Temporal Panel EBLUP	✓	✓	×	×	×	×
Unit-Level BHF Model	✓	✓	✓	×	×	×
OpenMP Parallel Bootstrap	✓	×	×	×	×	×
Non-Gaussian Likelihoods	6 Families	×	Box-Cox	Beta only	4 Families	3 Families
Deterministic Bayesian EBP	✓	×	×	×	×	✓
Exact Survey Dispersion Match	✓	×	×	✓	×	×
Modern Spatial Priors (BYM2/Leroux)	✓	×	×	Besag	Besag	BYM2
Two-Stage Hierarchical Calibration	✓	×	×	×	×	×
Official Statistics Reliability Flags	✓	×	×	×	×	×
Automated Production Mapping	✓	×	✓	×	×	✓
Multi-Sheet Dissemination Export	✓	×	✓	×	×	×

Table 1: Comparative landscape of Small Area Estimation packages in the R ecosystem.

suffers from divergent transitions on complex spatial graphs (Simpson, Rue, Riebler, Martins, and Sørbye 2017).

3. **The Institutional Benchmarking and Dissemination Gap:** In official statistics, model-based small area estimates must numerically harmonize with published higher-level totals (e.g., provincial and national direct estimates) to preserve institutional credibility and prevent contradictory policy decisions (You and Rao 2002; Wang, Fuller, and Qu 2008; Stefan and Hidirolou 2020). Existing packages rarely offer two-stage hierarchical benchmarking. Furthermore, software ecosystems typically end at model fitting, omitting essential dissemination steps such as automated reliability tagging based on official Relative Standard Error (RSE) thresholds, model concordance evaluation, and publication-ready spatial choropleth mapping.

As summarized in Table 1, the current software landscape is compartmentalized. Foundational toolkits such as **sae** (Molina and Marhuenda 2015) are implemented in interpreted R, lack multi-threading, and are restricted to Gaussian linear frameworks. Packages such as **emdi** (Kreutzmann, Pannier, Rojas-Perilla, Schmid, Templ, and Tzavidis 2019) offer visualizations but lack spatio-temporal dynamics and non-Gaussian spatial Bayesian likelihoods. Specialized packages such as **tipsae** (De Nicolò and Gardini 2022) implement Beta regression but depend on MCMC runtimes that scale poorly for panel graphs. Toolkits such as **mcmc-sae** (Boonstra 2022) remain simulation-based, while **SUMMER** (Li, Hsiao, Godwin, Martin, Wakefield, and Clark 2020) is tailored to demographic survey mortality rather than general survey design-variance matching.

To resolve these limitations, this paper introduces **fastsae**. Specifically, this paper makes four concrete, testable contributions:

1. **Compiled C++17 Core with OpenMP Acceleration:** We engineer an optimized linear algebra core in C++17 via **RcppArmadillo** that speeds up classical, spatial, and spatio-temporal Fay-Herriot estimation by up to 364× relative to **sae** and 12,300× relative to **emdi** at $D = 1,000$, while slashing peak memory consumption from gigabytes to under 25 megabytes.

2. **Unified Deterministic Bayesian Hierarchical Modeling via INLA with Exact Survey Matching:** We formulate a Hierarchical Bayes interface (`hb_area()`) spanning six probability distributions (Gaussian, Beta, Gamma, Binomial, Poisson, Negative Binomial) powered by Integrated Nested Laplace Approximations (Rue, Martino, and Chopin 2009). It embeds exact survey variance mappings—notably the Janicki (2020) Beta dispersion parameterization and Gamma CV matching—achieving $73\times$ to $544\times$ speedups over Stan-based MCMC while maintaining near-perfect inferential concordance ($r = 0.9993$).
3. **Two-Stage Hierarchical Benchmarking Engine:** We implement a modular calibration engine (`benchmark_sae()`) that simultaneously enforces consistency across multi-level administrative hierarchies (Level 0 National, Level 1 Provincial, Level 2 District) via ratio, difference, variance-weighted optimal, and logit-scale adjustments.
4. **End-to-End Official Statistics Dissemination Suite:** We provide turnkey post-estimation workflows, including automated reliability tagging (`add_reliability_flags()`) based on official NSO thresholds, model concordance diagnostics (`compare_sae()`), publication-ready choropleth mapping (`map_sae()`), and multi-sheet spreadsheet exports (`export_sae()`).

2. Methodological framework

2.1. Frequentist small area models

Classical Fay-Herriot model

Let m denote the number of small areas indexed by $i = 1, \dots, m$. The classical area-level model (Fay and Herriot 1979) is defined by:

$$y_i = \mathbf{x}_i^\top \boldsymbol{\beta} + u_i + e_i, \quad u_i \stackrel{\text{iid}}{\sim} \mathcal{N}(0, \sigma_u^2), \quad e_i \stackrel{\text{iid}}{\sim} \mathcal{N}(0, D_i), \quad (1)$$

where y_i is the design-based direct estimator of the target indicator θ_i , $\mathbf{x}_i$ is a $p \times 1$ vector of auxiliary covariates, D_i is the known design variance, and u_i is the area random effect. The Empirical Best Linear Unbiased Predictor (EBLUP) is given by:

$$\hat{\theta}_i = \hat{\gamma}_i y_i + (1 - \hat{\gamma}_i) \mathbf{x}_i^\top \hat{\boldsymbol{\beta}}, \quad \hat{\gamma}_i = \frac{\hat{\sigma}_u^2}{\hat{\sigma}_u^2 + D_i}, \quad (2)$$

where $\hat{\boldsymbol{\beta}} = (\mathbf{X}^\top \hat{\mathbf{V}}^{-1} \mathbf{X})^{-1} \mathbf{X}^\top \hat{\mathbf{V}}^{-1} \mathbf{y}$ with $\mathbf{V} = \text{diag}(\sigma_u^2 + D_1, \dots, \sigma_u^2 + D_m)$. Mean squared error is evaluated via the second-order Prasad-Rao approximation (Prasad and Rao 1990):

$$\text{MSE}(\hat{\theta}_i) \approx g_{1i}(\hat{\sigma}_u^2) + g_{2i}(\hat{\sigma}_u^2) + 2g_{3i}(\hat{\sigma}_u^2). \quad (3)$$

Spatial Fay-Herriot model

To accommodate spatial dependence across contiguous domains, Petrucci and Salvati (2006) and Pratesi and Salvati (2008) modeled the area random effects $\mathbf{u} = (u_1, \dots, u_m)^\top$ via a

Simultaneous Autoregressive (SAR) process:

$$\mathbf{u} = \rho \mathbf{W} \mathbf{u} + \boldsymbol{\epsilon}, \quad \boldsymbol{\epsilon} \sim \mathcal{N}_m(\mathbf{0}, \sigma_1^2 \mathbf{I}_m), \quad (4)$$

where ρ is the spatial autoregressive parameter and $\mathbf{W}$ is a row-standardized spatial proximity matrix. The marginal covariance matrix is $\mathbf{V}(\sigma_1^2, \rho) = \mathbf{G}(\sigma_1^2, \rho) + \mathbf{D}$, where $\mathbf{G}(\sigma_1^2, \rho) = \sigma_1^2 [(\mathbf{I}_m - \rho \mathbf{W})^\top (\mathbf{I}_m - \rho \mathbf{W})]^{-1}$. The spatial EBLUP is:

$$\hat{\boldsymbol{\theta}} = \mathbf{X} \hat{\boldsymbol{\beta}} + \hat{\mathbf{G}} \hat{\mathbf{V}}^{-1} (\mathbf{y} - \mathbf{X} \hat{\boldsymbol{\beta}}). \quad (5)$$

Variance components are estimated via profile REML over the valid spectral bounds $\rho \in (\lambda_{\min}^{-1}, \lambda_{\max}^{-1})$.

Spatio-temporal Fay-Herriot model

When small area data are collected across longitudinal waves $t = 1, \dots, T$, borrowing strength across space and time yields substantial efficiency gains (Marhuenda *et al.* 2013):

$$y_{it} = \mathbf{x}_{it}^\top \boldsymbol{\beta} + u_{1,it} + u_{2,it} + e_{it}, \quad (6)$$

where $\mathbf{u}_{1,t} = \rho_1 \mathbf{W} \mathbf{u}_{1,t} + \boldsymbol{\epsilon}_{1,t}$ represents a spatial SAR process and $u_{2,it} = \rho_2 u_{2,i,t-1} + \epsilon_{2,it}$ follows an AR(1) temporal process with $|\rho_2| < 1$. In **fastsae**, the block-structured covariance matrix is assembled and inverted using optimized Kronecker routines.

Unit-level Battese-Harter-Fuller model

When unit microdata are accessible alongside population auxiliary totals, **fastsae** implements the nested error regression model (Battese *et al.* 1988):

$$y_{dj} = \mathbf{x}_{dj}^\top \boldsymbol{\beta} + u_d + e_{dj}, \quad d = 1, \dots, D, \quad j = 1, \dots, n_d, \quad (7)$$

yielding EBLUP predictors $\hat{Y}_d = f_d \bar{y}_d + (1 - f_d)(\bar{\mathbf{X}}_{d,r}^\top \hat{\boldsymbol{\beta}} + \hat{\gamma}_d(\bar{y}_d - \bar{\mathbf{x}}_d^\top \hat{\boldsymbol{\beta}}))$, where $f_d = n_d/N_d$.

2.2. Bayesian non-Gaussian hierarchical EBP via INLA

Latent Gaussian models

For non-Gaussian indicators, **fastsae** formulates small area estimation within the Latent Gaussian Model (LGM) framework (Rue *et al.* 2009):

$$y_i \mid \eta_i, \boldsymbol{\theta} \sim \pi(y_i \mid \eta_i, \boldsymbol{\theta}), \quad \eta_i = g(\mu_i) = \mathbf{x}_i^\top \boldsymbol{\beta} + u_i + v_i, \quad (8)$$

where $g(\cdot)$ is a monotone link function, $v_i \stackrel{\text{iid}}{\sim} \mathcal{N}(0, \tau_v^{-1})$ is an unstructured effect, and $\mathbf{u} \sim \mathcal{N}(\mathbf{0}, \mathbf{Q}^{-1}(\boldsymbol{\theta}_s))$ is a spatial Gaussian Markov Random Field (GMRF). By default, **fastsae** utilizes the Full Laplace approximation (`strategy = "laplace"`) rather than the simplified Laplace expansion, guaranteeing high analytical fidelity even near boundary domains.

Exact survey dispersion and variance matching

To preserve design-based sampling properties in a Bayesian GLMM, **fastsae** injects distribution-specific dispersion parameters:

1. **Gaussian:** Precision is fixed to the inverse design variance: $\tau_i = 1/D_i$.
2. **Beta:** For proportions $y_i \in (0, 1)$ with logit link $\text{logit}(\mu_i) = \eta_i$, the Beta variance is $\text{Var}(y_i) = \mu_i(1 - \mu_i)/(1 + \phi_i)$. Equating this to design variance D_i yields the [Janicki \(2020\)](#) dispersion:

$$\phi_i = \max \left(\frac{y_i(1 - y_i)}{D_i} - 1, 1 \right). \quad (9)$$

3. **Gamma:** For positive skewed expenditures $y_i > 0$ with log link $\log(\mu_i) = \eta_i$, equating $\text{Var}(y_i) = \mu_i^2/s_i$ to D_i yields the shape precision parameter:

$$s_i = \frac{y_i^2}{D_i} = \frac{1}{\text{CV}_i^2}. \quad (10)$$

4. **Binomial, Poisson, and Negative Binomial:** Domain sample sizes n_i map to Ntrials , and population exposures E_i map to the log-link offset $\log(E_i)$.

Spatial priors and space-time interactions

The spatial random field $\mathbf{u}$ supports five graph topologies: Besag ICAR ([Besag, York, and Mollié 1991](#)), scaled BYM2 with Penalized Complexity (PC) priors ([Riebler, Sørbye, Simpson, and Rue 2016](#); [Simpson et al. 2017](#)), Leroux CAR with exact precision $\mathbf{Q} = \tau[(1 - \rho)\mathbf{I} + \rho(\text{diag}(\mathbf{d}) - \mathbf{W})]$ ([Leroux, Lei, and Breslow 2000](#)), and the Spatial Lag Model (SLM). For panel observations across $T \geq 2$, temporal trends support random walks (RW1, RW2) and AR(1) dynamics. Space-time interactions can be specified as additive, domain-specific, separable, or according to the four [Knorr-Held \(2000\)](#) interaction types ($\mathbf{Q}_{st} = \mathbf{Q}_s \otimes \mathbf{Q}_t$).

2.3. Two-stage hierarchical benchmarking calibration

In official statistics, model-based small area estimates $\hat{\theta}_{di}$ for domain i in province/region g must aggregate to officially published higher-level targets ([Rao and Molina 2015](#); [Stefan and Hidirolou 2020](#)). Let w_{di} denote domain population weights such that $W_g = \sum_{i \in g} w_{di}$ is the provincial total weight, and $W_{\text{nat}} = \sum_g W_g$ is the national total weight. Let T_g denote the provincial target and T_{nat} denote the national aggregate target.

fastsae implements a two-stage hierarchical benchmarking calibration:

1. **Stage 1: Harmonization of Provincial Targets to National Benchmark:** When initial provincial targets T_g (or initial aggregated model estimates $\hat{\theta}_g = \sum_{i \in g} (w_{di}/W_g) \hat{\theta}_{di}$) do not sum to the national target T_{nat} , calibrated provincial targets $\tilde{T}_g$ are computed to satisfy $\sum_g (W_g/W_{\text{nat}}) \tilde{T}_g = T_{\text{nat}}$.
2. **Stage 2: Domain-Level Calibration within Provinces:** Domain estimates within province g are calibrated to satisfy $\sum_{i \in g} (w_{di}/W_g) \hat{\theta}_{di}^{\text{BM}} = \tilde{T}_g$.

Across both stages, **fastsae** supports four calibration mechanisms:

- **Ratio Benchmarking** ([You and Rao 2002](#)): $\hat{\theta}_{di}^{\text{BM}} = \hat{\theta}_{di} \left(\tilde{T}_g / \sum_{j \in g} \tilde{w}_{dj} \hat{\theta}_{dj} \right)$.

- **Difference Benchmarking** (Rao and Molina 2015): $\hat{\theta}_{di}^{\text{BM}} = \hat{\theta}_{di} + \left(\tilde{T}_g - \sum_{j \in g} \tilde{w}_{dj} \hat{\theta}_{dj} \right)$.
- **Variance-Weighted Optimal Benchmarking** (Datta, Ghosh, Steorts, and Maples 2011; Steorts and Ghosh 2014): Adjustments are allocated inversely proportional to domain precision, minimizing quadratic estimation loss:

$$\hat{\theta}_{di}^{\text{BM}} = \hat{\theta}_{di} + \text{MSE}(\hat{\theta}_{di}) \left(\sum_{j \in g} \tilde{w}_{dj}^2 \text{MSE}(\hat{\theta}_{dj}) \right)^{-1} \left(\tilde{T}_g - \sum_{j \in g} \tilde{w}_{dj} \hat{\theta}_{dj} \right). \quad (11)$$

- **Logit-Scale Benchmarking**: For bounded proportions in $(0, 1)$, linear adjustments can push estimates outside $[0, 1]$. **fastsae** solves for an additive shift α_g on the logit scale:

$$\sum_{i \in g} \tilde{w}_{di} \text{logit}^{-1} \left(\text{logit}(\hat{\theta}_{di}) + \alpha_g \right) = \tilde{T}_g, \quad (12)$$

via robust one-dimensional root-finding, guaranteeing $\hat{\theta}_{di}^{\text{BM}} \in (0, 1)$.

2.4. Official statistics reliability classification

NSOs classify domain estimates into publication tiers based on the Relative Standard Error (RSE = $\frac{\sqrt{\text{MSE}(\hat{\theta}_i)}}{|\hat{\theta}_i|} \times 100\%$). **fastsae** provides `add_reliability_flags()` implementing institutional thresholding:

$$\text{Tier}_i = \begin{cases} \text{Reliable (Direct Publish),} & \text{RSE}_i < c_1, \\ \text{Use with Caution (Accompanying Warning),} & c_1 \leq \text{RSE}_i < c_2, \\ \text{Unreliable (Suppress or Aggregate),} & \text{RSE}_i \geq c_2, \end{cases} \quad (13)$$

where default thresholds $(c_1, c_2) = (20\%, 30\%)$ or $(25\%, 50\%)$ conform to standards adopted by agencies such as BPS-Statistics Indonesia and Eurostat.

3. Software design and implementation

The architecture of **fastsae** unites compiled linear algebra with high-level Bayesian GMRF modeling and production reporting. Figure 1 outlines the structural flow.

The frequentist engines (`eblup_fh()`, `eblup_sf()`, `eblup_stfh()`, `eblup_bhf()`) avoid temporary matrix allocations inside REML loops by reusing working buffers in-place. The bootstrap MSE loops are decorated with OpenMP pragmas (`#pragma omp parallel for`), delivering multi-core throughput scaling. The Bayesian engine (`hb_area()`) automates the translation of formulas, spatial adjacency matrices $\mathbf{W}$, and survey dispersion arguments into **R-INLA** calls. Standard S3 methods (`print()`, `summary()`, `plot()`, `coef()`, `fitted()`, `residuals()`) are implemented across all model classes.

4. Computational benchmarks and simulations

fastsae High-Level Software Architecture		
1. User-Facing Production API (R)		
Frequentist Modeling	Bayesian HB Engine	Production Workflow
eblup_fh()	hb_area(..., family = "beta")	benchmark_sae() (Two-Stage)
eblup_sfhh()	hb_area(..., family = "gamma")	compare_sae()
eblup_stfh()	hb_area(..., family = "poisson")	add_reliability_flags()
eblup_bhf()	hb_area(..., spatial = "bym2")	map_sae() / export_sae()
2. Core Computational Backends		
Compiled C++17 Core	INLA Deterministic Engine	Spatial Linear Algebra
RcppArmadillo Matrix Algebra	Sparse Cholesky Factorization	Row-normalized Proximity
OpenMP Threaded Bootstrap	Full Laplace Approximation	PC Priors & GMRF Graphs
In-Place Fisher-Scoring Updates	Survey Dispersion Injections	Spectral Eigenvalue Bounds

Figure 1: Structural architecture and functional components of **fastsae**.

4.1. Frequentist scalability and memory efficiency: **fastsae** vs. **sae** and **emdi**

We evaluated computational speed and peak memory consumption across domain sizes $D \in \{30, 50, 100, 250, 500, 1000\}$ using simulated datasets with 5 auxiliary covariates. Benchmarks were conducted using the **bench** package on an Apple Silicon processor under R 4.4.2.

Table 2 demonstrates the dramatic scalability gains of **fastsae**. At $D = 1,000$, classical FH completes in 4 milliseconds in **fastsae** compared to 1.5 seconds in **sae** ($364\times$ slower) and 51 seconds in **emdi** ($12,343\times$ slower). For spatio-temporal models at $D = 1,000$, **sae** required nearly 30 minutes (1,786 seconds) and consumed 32.5 gigabytes of RAM, whereas **fastsae** finished in 58 seconds using under 1 megabyte of memory. **Operational Takeaway:** **fastsae** eliminates out-of-memory bottlenecks, enabling routine spatio-temporal estimation across thousands of domains on standard laptop hardware.

4.2. Bayesian Beta Hierarchical Modeling scalability: **fastsae** vs. **tipsae**

For bounded rates, we benchmarked `hb_area(..., family = "beta")` against **tipsae** (De Nicolò and Gardini 2022) using Stan (4 chains, 2,000 iterations).

As shown in Table 3, deterministic Laplace approximations in **fastsae** scale flatly with domain count: fitting a spatial Beta model at $D = 1,000$ requires only 1.73 seconds in **fastsae** compared to 404 seconds in **tipsae** ($234\times$ speedup). In spatio-temporal panel models ($D = 50, T = 5$), **tipsae** required 15.1 minutes (908 seconds), whereas **fastsae** completed in 1.67 seconds ($544\times$ speedup). Despite this speed differential, inferential concordance on $D = 1,000$ yields Pearson correlation $r = 0.9993$, Spearman $\rho = 0.9992$, and Mean Absolute

D	Execution Time: Median (seconds)			Peak Memory Footprint (MB)		
	fastsae	sae	emdi	fastsae	sae	emdi
<i>Fay-Herriot EBLUP (Classical)</i>						
30	0.00059	0.00155	0.01088	0.005	0.153	4.18
100	0.00055	0.00344	0.05051	0.016	0.918	37.34
500	0.00140	0.19673	5.86366	0.076	18.239	880.31
1000	0.00413	1.50363	51.0224	0.186	71.977	3,773.60
<i>Spatial Fay-Herriot (SEBLUP)</i>						
30	0.00073	0.00410	0.01071	0.030	1.373	4.18
100	0.00473	0.07054	0.05299	0.259	17.976	37.34
500	0.14205	7.11793	3.13662	5.857	776.407	949.86
1000	0.83216	67.1518	51.0593	23.155	1,920.01	3,772.57
<i>Spatio-Temporal Fay-Herriot (STEBLUP, $T = 5$)</i>						
30	0.03674	0.11856	–	0.027	43.801	–
100	2.29928	17.9425	–	0.090	2,551.10	–
500	10.3596	253.343	–	0.448	8,124.37	–
1000	58.1005	1,786.49	–	0.894	32,478.3	–

Table 2: Frequentist runtime and peak memory benchmark across domain sample sizes (D). **fastsae** avoids quadratic RAM growth and achieves up to $364\times$ speedup for FH, $80\times$ for SFH, and $31\times$ for ST-FH at $D = 1,000$.

D	Cross-Sectional Beta		Spatial Beta (Besag)		Spatio-Temporal Beta	
	fastsae	tipsae	fastsae	tipsae	fastsae	tipsae
30	1.24 s	2.35 s	0.61 s	25.72 s	1.44 s	540.57 s
50	1.33 s	3.77 s	0.62 s	65.72 s	1.67 s	908.17 s
100	1.27 s	8.24 s	1.37 s	100.02 s	–	–
250	1.31 s	26.28 s	1.32 s	155.90 s	–	–
500	1.41 s	64.21 s	1.43 s	248.82 s	–	–
1000	1.64 s	187.16 s	1.73 s	404.34 s	–	–

Table 3: Execution time comparison for Bayesian Beta models. Point estimates between **fastsae** and **tipsae** achieve $r = 0.9993$ concordance.

Difference MAD = 0.0197. **Operational Takeaway:** Deterministic Laplace approximation in **fastsae** achieves inferential equivalence with fully converged Hamiltonian MCMC in seconds rather than minutes or hours.

4.3. Monte Carlo finite population validation against ground truth

To assess frequentist design properties, we conducted design-based Monte Carlo simulations under known ground truth at national scale ($D = 500$ small areas, $R = 25$ independent sample draws) using a synthetic finite population of $N = 250,000$ individuals.

Table 4 demonstrates that **hb_area()** reduces empirical RMSE from 0.0719 (direct survey estimator) to 0.0481, yielding a $1.92\times$ empirical efficiency gain. Relative Root Mean Squared Error drops from 27.1% to 19.1%. Comparing Bayesian engines, **fastsae** achieves a slightly lower empirical RMSE than **tipsae** (0.0481 vs. 0.0487) and produces narrower credible inter-

Estimator	Emp. RMSE	RRMSE (%)	MARB (%)	95% CrI Cov.	CrI Width	Runtime
Direct Sample Estimator	0.0719	27.1%	21.7%	—	—	—
fastsae (Spatial Besag)	0.0481	19.1%	16.9%	89.0%	0.173	4.17 s
tipsae (Spatial Besag)	0.0487	18.7%	15.7%	93.7%	0.202	36.86 s

Table 4: Finite population Monte Carlo simulation results ($D = 500$, $R = 25$). **fastsae** achieves an efficiency gain of $1.92\times$ over the direct survey estimator.

vals (0.173 vs. 0.202) while requiring nearly $9\times$ less runtime per fit. **Operational Takeaway:** **fastsae** produces well-calibrated, highly efficient area-level predictions that closely track true finite-population parameters while scaling effortlessly to national-scale administrative systems.

5. Applied illustration and production workflow

We demonstrate an end-to-end NSO production workflow using the benchmark dataset `sim_area` ($D = 42$ domains across 6 provincial groups).

5.1. Loading data, spatial weights, and fitting models

```
R> library("fastsae")
R> data("sim_area")
R> sim_area$prov <- rep(paste0("Prov_", 1:6), each = 7)
R> sim_area$pop <- round(runif(nrow(sim_area), 15000, 45000))
R> # Construct row-standardized spatial proximity matrix
R> coords <- cbind(sim_area$x_coord, sim_area$y_coord)
R> dists <- as.matrix(dist(coords))
R> W <- (dists < quantile(dists, 0.18)) * 1
R> diag(W) <- 0
R> rs <- rowSums(W); rs[rs == 0] <- 1; W_std <- W / rs
R> # Fit Frequentist and Bayesian Models
R> fit_fh <- eblup_fh(y_gaussian ~ x1 + x2, vardir = "varDir",
+   data = sim_area, method = "REML", print_result = FALSE)
R> fit_sfh <- eblup_sfh(y_gaussian ~ x1 + x2, vardir = "varDir",
+   domain = "area", W = W_std, data = sim_area,
+   method = "REML", print_result = FALSE)
R> fit_beta <- hb_area(y_beta ~ x1 + x2, vardir = "varDir",
+   data = sim_area, family = "beta", spatial = "bym2",
+   W = W, print_result = FALSE)
```

5.2. Automated model concordance evaluation

Practitioners can immediately compare competitive models via `compare_sae()`:

```
R> comp <- compare_sae(fit_fh, fit_sfh, names = c("FH", "SFH"))
R> comp$metrics[c(1, 2, 4, 6, 8, 9), ]
```

	Metric	Value
1	Matched Domains	42
2	Pearson Correlation (r)	0.9993
4	Mean Absolute Difference (MAE)	0.0357
6	Mean MSE Ratio (FH / SFH)	1.0035
8	Domains where SFH has lower MSE	8 / 42
9	Percentage of domains where SFH is more efficient (%)	19.0%

5.3. Two-stage hierarchical benchmarking calibration

To harmonize domain estimates with an official national poverty target ($T_{\text{nat}} = 0.50$), we apply two-stage calibration:

```
R> bm <- benchmark_sae(fit_fh, national_target = 0.50, group = "prov",
+   weight = "pop", method = "ratio", type = "mean")
R> head(bm[, c("domain", "group", "original", "benchmarked", "rel_adjustment_pct")], 4)
```

```
=== fastsae Small Area Benchmark Calibration ===
Method: Ratio (Multiplicative Adjustment)
Target Type: Weighted Mean / Rate
Total Domains: 4
```

```
First 6 benchmarked domains (Level 2):
  domain group original benchmarked rel_adjustment_pct
1      1 Prov_1  2.6646558   0.8665880          -67.47843
2      2 Prov_1 -1.6547233  -0.5381420          -67.47843
3      3 Prov_1  1.2504674   0.4066717          -67.47843
4      4 Prov_1 -0.7988192  -0.2597886          -67.47843
```

The calibrated estimates satisfy $\sum(w_i / \sum w_i) \hat{\theta}_i^{\text{BM}} = 0.50$ exactly.

5.4. Official statistics reliability classification

Using `add_reliability_flags()`, domains are classified into official dissemination tiers:

```
R> flags <- add_reliability_flags(fit_beta, thresholds = c(20, 30))
R> table(flags$reliability)
```

```
Reliable (< 20%) Use with Caution (20-30%)
              31                      8
Unreliable (≥ 30%)
              3
```

5.5. Model diagnostics and visualization

Diagnostics via `diagnose()` and `autoplot()` provide visual quality verification:

```
R> diag_beta <- diagnose(fit_beta)
R> print(autoplot(diag_beta))
```




Figure 2: Model diagnostic panels generated via `autoplot.fastsaе_diagnose()`.

6. Summary and discussion

Small area estimation in official statistics requires reconciling mathematical rigor with strict computational and operational constraints. **fastsaе** resolves these challenges by bridging high-performance compiled linear algebra, deterministic Bayesian approximations, and production-ready dissemination tools.

As synthesized in Table 5, **fastsaе** provides substantial efficiency gains across both frequentist and Bayesian paradigms. Future developments will focus on expanding unit-level spatial EBP under non-Gaussian likelihoods, integrating zero-inflated Beta (ZIB) specifications for extreme sparse indicators, and supporting direct spatial polygon streaming.

Computational details

The results in this paper were obtained using R 4.4.2 with **fastsaе** 0.1.0, **R-INLA** 24.06.27, **sae** 1.3, **emdi** 2.2.1, **tipsae** 1.0.1, and **ggplot2** 3.5.1. R and packages from CRAN are available at <https://CRAN.R-project.org/>. **R-INLA** is available from <https://www.r-inla.org/>. The **fastsaе** package is available on GitHub at <https://github.com/ridsonap/fastsaе>.

Acknowledgments

The authors thank the Department of Computational Statistics and Department of Statistics at Politeknik Statistika STIS for institutional support. We also acknowledge the valuable domain feedback provided by colleagues at BPS-Statistics Indonesia.

Dimension	Observed Empirical Evidence	Institutional Implication (“So What”)
Frequentist Scalability	364× speedup for FH and 80× for SFH; RAM under 25 MB at $D = 1,000$ (vs 32 GB in <i>sae</i>).	Eliminates server memory exhaustion, allowing desktop computers to execute nationwide bootstrap MSE routines.
Bayesian via INLA	73× to 544× speedup over Stan MCMC; point correlation $r = 0.9993$ on $D = 1,000$.	Eliminates manual MCMC chain diagnostics and divergent transition failures in routine production timelines.
Design-Based Ground Truth	1.92× empirical efficiency gain over direct estimator; 89% empirical coverage.	Delivers substantial MSE reductions while maintaining nominal interval coverage under repeated sampling.
Hierarchical Benchmarking	Simultaneous Level 0 and Level 1 target alignment with logit boundary safety.	Guarantees numeric consistency with published aggregates, eliminating public confusion without violating probability bounds.
Turnkey Dissemination	Direct RSE classification, spatial choropleths, and multi-sheet exports.	Collapses multi-script analyst workflows into standardized, reproducible publication pipelines.

Table 5: Synthesis of empirical findings and operational implications of **fastsae**.

References

- Battese GE, Harter RM, Fuller WA (1988). “An Error-Components Model for Prediction of County Crop Areas Using Survey and Satellite Data.” *Journal of the American Statistical Association*, **83**(401), 28–36. doi:10.1080/01621459.1988.10478561.
- Besag J, York J, Mollié A (1991). “Bayesian Image Restoration, with Two Applications in Spatial Statistics.” *Annals of the Institute of Statistical Mathematics*, **43**(1), 1–20. doi:10.1007/BF00116466.
- Boonstra HJ (2022). *mcmcsae: Markov Chain Monte Carlo Small Area Estimation*. R package version 0.7.3, URL <https://CRAN.R-project.org/package=mcmcsae>.
- Boubeta M, Lombardía MJ, Morales D (2016). “Empirical Best Prediction Under Area-Level Poisson Mixed Models.” *TEST*, **25**(3), 540–569. doi:10.1007/s11749-015-0468-x.
- Burgard JP, Krause J, Merkle R (2020). “Computational Aspects of Small Area Estimation.” *Statistical Papers*, **61**(6), 2533–2555. doi:10.1007/s00362-020-01185-w.
- Datta GS, Ghosh M, Steorts R, Maples J (2011). “Bayesian Benchmarking with Applications to Small Area Estimation.” *TEST*, **20**(3), 574–588. doi:10.1007/s11749-010-0218-y.
- De Nicolò S, Gardini A (2022). “Mapping Proportions and Indicators on the Unit Interval in Small Area Estimation with the R Package **tipsae**.” *The R Journal*, **14**(4), 166–184. doi:10.32614/RJ-2022-049.

- Fabrizi E, Ferrante MR, Trivisano C (2011). “Hierarchical Bayes Approaches to Model-Based Small Area Estimation with Application to the Estimation of Poverty Indicators.” *Computational Statistics & Data Analysis*, **55**(4), 1636–1647. doi:[10.1016/j.csda.2010.10.015](https://doi.org/10.1016/j.csda.2010.10.015).
- Fay RE, Herriot RA (1979). “Estimates of Income for Small Places: An Application of James-Stein Procedures to Census Data.” *Journal of the American Statistical Association*, **74**(366a), 269–277. doi:[10.1080/01621459.1979.10482505](https://doi.org/10.1080/01621459.1979.10482505).
- González-Manteiga W, Lombardía MJ, Molina I, Morales D, Santamaría L (2008). “Bootstrap Mean Squared Error of a Small-Area EBLUP.” *Journal of Statistical Computation and Simulation*, **78**(5), 443–462. doi:[10.1080/00949650601141811](https://doi.org/10.1080/00949650601141811).
- Janicki MW (2020). “Properties and Applications of Beta Regression Models for Small Area Estimation.” *Technical Report Research Report Series (Statistics #2020-02)*, Center for Statistical Research and Methodology, U.S. Census Bureau, Washington, D.C.
- Knorr-Held L (2000). “Bayesian Modelling of Spatio-Temporal Disease Mapping.” *Statistics in Medicine*, **19**(17–18), 2555–2567.
- Kreutzmann AK, Pannier S, Rojas-Perilla N, Schmid T, Templ M, Tzavidis N (2019). “The R Package **emdi** for Estimating and Mapping Indicators to Monitor Poverty and Inequality.” *Journal of Statistical Software*, **91**(7), 1–33. doi:[10.18637/jss.v091.i07](https://doi.org/10.18637/jss.v091.i07).
- Leroux BG, Lei X, Breslow N (2000). “Estimation of Disease Rates in Small Areas: A New Mixed Model for Spatial Dependence.” In *Statistical Models in Epidemiology, the Environment, and Clinical Trials*, pp. 135–178. Springer-Verlag, New York. doi:[10.1007/978-1-4612-1284-3_4](https://doi.org/10.1007/978-1-4612-1284-3_4).
- Li ZR, Hsiao Y, Godwin J, Martin BD, Wakefield J, Clark SJ (2020). “Changes in the Spatial Distribution of the Under-Five Mortality Rate: Small-Area Analysis of 122 DHS Surveys in 262 Subregions of 35 Countries in Africa.” *PLOS Medicine*, **17**(1), e1003061. doi:[10.1371/journal.pmed.1003061](https://doi.org/10.1371/journal.pmed.1003061).
- Marhuenda Y, Molina I, Morales D (2013). “Small Area Estimation with Spatio-Temporal Fay-Herriot Models.” *Computational Statistics & Data Analysis*, **58**, 308–325. doi:[10.1016/j.csda.2012.09.002](https://doi.org/10.1016/j.csda.2012.09.002).
- Molina I, Marhuenda Y (2015). “**sae**: An R Package for Small Area Estimation.” *The R Journal*, **7**(1), 73–87. doi:[10.32614/RJ-2015-007](https://doi.org/10.32614/RJ-2015-007).
- Petrucchi A, Salvati N (2006). “Small Area Estimation for Spatial Data in Presence of Correlated Random Area Effects.” *Journal of Official Statistics*, **22**(3), 407–428.
- Pfeffermann D (2013). “New Important Developments in Small Area Estimation.” *Statistical Science*, **28**(1), 40–68. doi:[10.1214/12-STS395](https://doi.org/10.1214/12-STS395).
- Prasad NGN, Rao JNK (1990). “The Estimation of the Mean Squared Error of Small-Area Estimators.” *Journal of the American Statistical Association*, **85**(409), 163–171. doi:[10.1080/01621459.1990.10475320](https://doi.org/10.1080/01621459.1990.10475320).

- Pratesi M (ed.) (2016). *Analysis of Poverty Data by Small Area Estimation*. John Wiley & Sons, Chichester, United Kingdom. doi:[10.1002/9781118814963](https://doi.org/10.1002/9781118814963).
- Pratesi M, Salvati N (2008). “Small Area Estimation: The EBLUP Estimator Based on Spatially Correlated Random Area Effects.” *Statistical Methods & Applications*, **17**(1), 113–141. doi:[10.1007/s10260-007-0061-9](https://doi.org/10.1007/s10260-007-0061-9).
- Rao JNK, Molina I (2015). *Small Area Estimation*. 2nd edition. John Wiley & Sons, Hoboken, New Jersey. doi:[10.1002/9781118735855](https://doi.org/10.1002/9781118735855).
- Riebler A, Sørbye SH, Simpson D, Rue H (2016). “An Intuitive Bayesian Spatial Model for Disease Mapping that Accounts for Scaling.” *Statistical Methods in Medical Research*, **25**(4), 1145–1165. doi:[10.1177/0962280216660421](https://doi.org/10.1177/0962280216660421).
- Rue H, Martino S, Chopin N (2009). “Approximate Bayesian Inference for Latent Gaussian Models by Using Integrated Nested Laplace Approximations.” *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, **71**(2), 319–392. doi:[10.1111/j.1467-9868.2008.00700.x](https://doi.org/10.1111/j.1467-9868.2008.00700.x).
- Simpson D, Rue H, Riebler A, Martins TG, Sørbye SH (2017). “Penalising Model Component Complexity: A Principled, Practical Approach to Constructing Priors.” *Statistical Science*, **32**(1), 1–28. doi:[10.1214/16-ST576](https://doi.org/10.1214/16-ST576).
- Stefan M, Hidiroglou MA (2020). “Benchmarking Small Area Estimators with a Two-Stage Calibration Approach.” *Survey Methodology*, **46**(2), 241–262.
- Steorts RC, Ghosh M (2014). “On Benchmarking with Applications to the Estimation of Mean Squared Errors of Small Area Estimators.” *Biometrika*, **101**(2), 397–412. doi:[10.1093/biomet/ast053](https://doi.org/10.1093/biomet/ast053).
- Wang J, Fuller WA, Qu Y (2008). “Small Area Estimation Under a Restriction.” *Survey Methodology*, **34**(1), 29–36.
- You Y, Rao JNK (2002). “Small Area Estimation Using Unmatched Sampling and Linking Models.” *The Canadian Journal of Statistics*, **30**(1), 3–15. doi:[10.2307/3315862](https://doi.org/10.2307/3315862).

Affiliation:

Ridson Al Farizal P.
Department of Computational Statistics
Politeknik Statistika STIS
Jl. Otto Iskandardinata No. 64C
Jakarta Timur 13330, Indonesia
E-mail: alfrzlp@gmail.com

and

Azka Ubaidillah
Department of Statistics
Politeknik Statistika STIS
Jl. Otto Iskandardinata No. 64C
Jakarta Timur 13330, Indonesia
E-mail: azka@stis.ac.id