Fastsae: High-Performance Frequentist and Bayesian Small Area Estimation in R

by Ridson Al Farizal P. and Azka Ubaidillah

Abstract Small area estimation (SAE) plays a pivotal operational role in official statistics, enabling the production of reliable disaggregated socioeconomic, health, and demographic indicators when domain-specific sample sizes are small or zero. However, practitioners in national statistical offices (NSOs) face a triple computational and operational barrier: frequentist spatial and spatio-temporal models with bootstrap mean squared error (MSE) estimation suffer from severe memory bottlenecks and slow runtimes in interpreted languages; non-Gaussian indicators (such as rates, skewed expenditures, and counts) traditionally require computationally demanding Markov chain Monte Carlo (MCMC) algorithms that impede routine production; and model estimates often fail to aggregate to published higher-level totals, violating institutional consistency. We introduce fastsae, an R package engineered for high-performance frequentist and Bayesian small area estimation. On the frequentist side, fastsae implements optimized C++17 routines via RcppArmadillo paired with OpenMP multi-threading, accelerating the classical Fay-Herriot, spatial Fay-Herriot, spatio-temporal Fay-Herriot, and unit-level Battese-Harter-Fuller models alongside their analytical and bootstrap MSE estimators. On the Bayesian side, fastsae provides a unified empirical best prediction (EBP) interface across six distribution families (Gaussian, Beta, Gamma, Binomial, Poisson, Negative Binomial), fitted by Integrated Nested Laplace Approximations (INLA) by default or by MCMC in Python (JAX/NumPyro) through interchangeable engine objects. Crucially, fastsae introduces exact survey dispersion mappings that respect design-based direct sampling variances and accommodates modern spatial priors (BYM, scaled BYM2, Leroux CAR, and SLM) as well as Knorr-Held space-time interaction dynamics. To bridge estimation and dissemination, the package features a two-stage hierarchical benchmarking engine for simultaneous provincial and national calibration, an automated reliability classification engine aligned with NSO standards, multi-model concordances, automated ggplot2 spatial mapping, and multi-sheet spreadsheet exports. Extensive benchmarks and finite population Monte Carlo simulations demonstrate large computational speedups at 1,000 domains (1,263x for Fay-Herriot and 67x for spatial Fay-Herriot relative to sae, 217x for spatial Beta relative to Stan-based MCMC) and reductions of more than 99% in allocated memory, without loss of inferential precision.

1 Introduction

In modern evidence-based governance, policy interventions are increasingly decentralized. Monitoring the United Nations Sustainable Development Goals (SDGs), allocating municipal fiscal transfers, targeting social safety net subsidies, and responding to localized disease outbreaks all require dependable statistical indicators at highly disaggregated geographic administrative levels (Rao and Molina, 2015; Pfeffermann, 2013). However, national statistical offices (NSOs) face structural sample size limitations. Large-scale household surveys—such as the American Community Survey (ACS) in the United States, the European Union Statistics on Income and Living Conditions (EU-SILC), and the National Socioeconomic Survey (Susenas) in Indonesia—are powered to yield acceptable sampling errors only at national or broad regional levels. When partitioned into sub-domains (e.g., municipalities, districts, or health zones), domain-specific sample sizes become sparse or zero, causing direct survey estimators to suffer from high sampling variance and unacceptable coefficients of variation (Burgard et al., 2020; Pratesi, 2016).

Small area estimation (SAE) overcomes these limitations by "borrowing strength'' across domains and time periods via linking models that combine survey responses with auxiliary administrative registers, census aggregates, or geospatial layers (Rao and Molina, 2015). The cornerstone of this literature remains the area-level Fay-Herriot (FH) model (Fay and Herriot, 1979), extended to account for spatial dependencies via Simultaneous Autoregressive (SAR) structures (Petrucci and Salvati, 2006; Pratesi and Salvati, 2008) and temporal dynamics via panel autoregressive processes (Marhuenda et al., 2013). In the presence of micro-level auxiliary data, the nested error regression model of Battese et al. (1988) (BHF) provides unit-level empirical best linear unbiased predictors (EBLUP).

Despite these theoretical advances, operationalizing small area models within production environments in NSOs has been hindered by three structural bottlenecks: %

1. **Frequentist Computational and Memory Bottlenecks**: Existing implementations of spatial and spatio-temporal Fay-Herriot models rely on interpreted languages. Estimating model variance components via Restricted Maximum Likelihood (REML) requires repeated inversions of $D \times D$ or $DT \times DT$ covariance matrices. When evaluating Mean Squared Error (MSE)

Feature / Package	fastsae	sae	emdi	tipsae	mcmcsae	SUMMER
Computational Engine	C++17 / INLA	Pure R	R / nlme	Stan HMC	C++ / MCMC	INLA
Frequentist EBLUP (FH/SFH)	✓	✓	✓	×	×	×
Spatio-Temporal Panel EBLUP	✓	✓	×	×	×	×
Unit-Level BHF Model	✓	✓	✓	×	×	×
OpenMP Parallel Bootstrap	✓	×	×	×	×	×
Non-Gaussian Likelihoods	6 Families	×	Box-Cox	Beta only	4 Families	3 Families
Deterministic Bayesian EBP	✓	×	×	×	×	✓
Exact Survey Dispersion Match	✓	×	×	✓	×	×
Modern Spatial Priors (BYM2/Leroux)	✓	×	×	Besag	Besag	BYM2
Two-Stage Hierarchical Calibration	✓	×	×	×	×	×
Official Statistics Reliability Flags	✓	×	×	×	×	×
Automated Production Mapping	✓	×	✓	×	×	✓
Multi-Sheet Dissemination Export	✓	×	✓	×	×	×

Table 1: Comparative landscape of Small Area Estimation packages in the R ecosystem.

through parametric or non-parametric bootstrap resampling (González-Manteiga et al., 2008), runtimes escalate to hours. Crucially, memory allocation for dense covariance arrays grows quadratically ($O(D^2)$ to $O(D^2T^2)$), triggering out-of-memory (OOM) crashes on large national grids ($D \geq 500$).

2. **Distributional Inflexibility and MCMC Friction in Bayesian SAE**: Real-world welfare indicators are rarely Gaussian. Rates and proportions (poverty headcounts, unemployment, stunting) are strictly bounded in $(0, 1)$, while expenditures exhibit severe positive skewness (Gamma), and demographic events represent discrete counts (Poisson, Negative Binomial) (Janicki, 2020; Fabrizi et al., 2011; Boubeta et al., 2016). Under non-Gaussian likelihoods, linear EBLUP is invalid, requiring Empirical Best Prediction (EBP) within a Hierarchical Bayesian (HB) framework. Currently, Bayesian SAE software relies on Markov Chain Monte Carlo (MCMC) sampling via **Stan** or **JAGS** (De Nicolò and Gardini, 2022). MCMC sampling requires lengthy runtimes, complex convergence diagnostics ($\hat{R}$, trace plots), and often suffers from divergent transitions on complex spatial graphs (Simpson et al., 2017).
3. **The Institutional Benchmarking and Dissemination Gap**: In official statistics, model-based small area estimates must numerically harmonize with published higher-level totals (e.g., provincial and national direct estimates) to preserve institutional credibility and prevent contradictory policy decisions (You and Rao, 2002; Wang et al., 2008; Stefan and Hidiroglou, 2020). Existing packages rarely offer two-stage hierarchical benchmarking. Furthermore, software ecosystems typically end at model fitting, omitting essential dissemination steps such as automated reliability tagging based on official Relative Standard Error (RSE) thresholds, model concordance evaluation, and publication-ready spatial choropleth mapping.

As summarized in Table 1, the current software landscape is compartmentalized. Foundational toolkits such as **sae** (Molina and Marhuenda, 2015) are implemented in interpreted R, lack multithreading, and are restricted to Gaussian linear frameworks. Packages such as **emdi** (Kreutzmann et al., 2019) offer visualizations but lack spatio-temporal dynamics and non-Gaussian spatial Bayesian likelihoods. Specialized packages such as **tipsae** (De Nicolò and Gardini, 2022) implement Beta regression but depend on MCMC runtimes that scale poorly for panel graphs. Toolkits such as **mcmcsae** (Boonstra, 2022) remain simulation-based, while **SUMMER** (Li et al., 2020) is tailored to demographic survey mortality rather than general survey design-variance matching.

To resolve these limitations, this paper introduces **fastsae**. Specifically, this paper makes four concrete, testable contributions: %

1. **Compiled C++17 Core with OpenMP Acceleration**: We engineer an optimized linear algebra core in C++17 via **RcppArmadillo** that speeds up classical, spatial, and spatio-temporal Fay-Herriot estimation with analytical MSE by $1,263\times$, $67\times$, and $24\times$ relative to **sae** (and the classical model by $42,874\times$ relative to **emdi**) at $D = 1,000$, while reducing allocated memory from gigabytes to a few megabytes. Sparse Cholesky factorizations and structured Woodbury identities avoid every dense $D \times D$ (or $DT \times DT$) inverse.
2. **Unified Deterministic Bayesian Hierarchical Modeling via INLA with Exact Survey Matching**: We formulate a Hierarchical Bayes interface (`hb_area()`) spanning six probability distributions (Gaussian, Beta, Gamma, Binomial, Poisson, Negative Binomial) powered by Integrated Nested Laplace Approximations (Rue et al., 2009), with an optional MCMC engine in Python (JAX/NumPyro) selected through the same interface (`engine = inla_engine()` or `engine = jax_engine()`). It embeds exact survey variance mappings—notably the Janicki (2020) Beta dispersion parameterization and Gamma CV matching—achieving $117\times$ to $530\times$ speedups over

Stan-based MCMC at the largest benchmarked sizes while maintaining near-perfect inferential concordance ($r = 0.9993$).

3. **Two-Stage Hierarchical Benchmarking Engine**: We implement a modular calibration engine (`benchmark_sae()`) that simultaneously enforces consistency across multi-level administrative hierarchies (Level 0 National, Level 1 Provincial, Level 2 District) via ratio, difference, variance-weighted optimal, and logit-scale adjustments.
4. **End-to-End Official Statistics Dissemination Suite**: We provide turnkey post-estimation workflows, including automated reliability tagging (`add_reliability_flags()`) based on official NSO thresholds, model concordance diagnostics (`compare_sae()`), publication-ready choropleth mapping (`map_sae()`), and multi-sheet spreadsheet exports (`export_sae()`).

%% – Methodological Framework ————————————————————-

2 Methodological framework

2.1 Frequentist small area models

Classical Fay-Herriot model

Let m denote the number of small areas indexed by $i = 1, \ldots, m$. The classical area-level model (Fay and Herriot, 1979) is defined by: %

$$y_i = \boldsymbol{x}_i^\top \boldsymbol{\beta} + u_i + e_i, \quad u_i \overset{\text{iid}}{\sim} \mathcal{N}(0, \sigma_u^2), \quad e_i \overset{\text{ind}}{\sim} \mathcal{N}(0, D_i), \tag{1}$$

% where y_i is the design-based direct estimator of the target indicator θ_i, $\boldsymbol{x}_i$ is a $p \times 1$ vector of auxiliary covariates, D_i is the known design variance, and u_i is the area random effect. The Empirical Best Linear Unbiased Predictor (EBLUP) is given by: %

$$\hat{\theta}_i = \hat{\gamma}_i y_i + (1 - \hat{\gamma}_i) \boldsymbol{x}_i^\top \hat{\boldsymbol{\beta}}, \quad \hat{\gamma}_i = \frac{\hat{\sigma}_u^2}{\hat{\sigma}_u^2 + D_i}, \tag{2}$$

% where $\hat{\boldsymbol{\beta}} = (\boldsymbol{X}^\top \hat{\boldsymbol{V}}^{-1} \boldsymbol{X})^{-1} \boldsymbol{X}^\top \hat{\boldsymbol{V}}^{-1} \boldsymbol{y}$ with $\boldsymbol{V} = \text{diag}(\sigma_u^2 + D_1, \ldots, \sigma_u^2 + D_m)$. Mean squared error is evaluated via the second-order Prasad-Rao approximation (Prasad and Rao, 1990): %

$$\text{MSE}(\hat{\theta}_i) \approx g_{1i}(\hat{\sigma}_u^2) + g_{2i}(\hat{\sigma}_u^2) + 2g_{3i}(\hat{\sigma}_u^2). \tag{3}$$

Spatial Fay-Herriot model

To accommodate spatial dependence across contiguous domains, Petrucci and Salvati (2006) and Pratesi and Salvati (2008) modeled the area random effects $\boldsymbol{u} = (u_1, \ldots, u_m)^\top$ via a Simultaneous Autoregressive (SAR) process: %

$$\boldsymbol{u} = \rho \boldsymbol{W} \boldsymbol{u} + \boldsymbol{\epsilon}, \quad \boldsymbol{\epsilon} \sim \mathcal{N}_m(\boldsymbol{0}, \sigma_1^2 \boldsymbol{I}_m), \tag{4}$$

% where ρ is the spatial autoregressive parameter and $\boldsymbol{W}$ is a row-standardized spatial proximity matrix. The marginal covariance matrix is $\boldsymbol{V}(\sigma_1^2, \rho) = \boldsymbol{G}(\sigma_1^2, \rho) + \boldsymbol{D}$, where $\boldsymbol{G}(\sigma_1^2, \rho) = \sigma_1^2 [(\boldsymbol{I}_m - \rho \boldsymbol{W})^\top (\boldsymbol{I}_m - \rho \boldsymbol{W})]^{-1}$. The spatial EBLUP is: %

$$\hat{\boldsymbol{\theta}} = \boldsymbol{X} \hat{\boldsymbol{\beta}} + \hat{\boldsymbol{G}} \hat{\boldsymbol{V}}^{-1} (\boldsymbol{y} - \boldsymbol{X} \hat{\boldsymbol{\beta}}). \tag{5}$$

% The SAR process is defined on all domains of $\boldsymbol{W}$, including unsampled ones, which are predicted by spatial kriging. Variance components (σ_1^2, ρ) are estimated by REML or ML Fisher scoring with $|\rho| \leq 0.999$, falling back to a constrained (projected) Fisher scoring when the classical iteration ends at a negative variance. All quantities are obtained from sparse Cholesky factorizations of $\boldsymbol{A} = (\boldsymbol{I} - \rho \boldsymbol{W})^\top (\boldsymbol{I} - \rho \boldsymbol{W})$ and $\boldsymbol{A} + \sigma_1^2 \boldsymbol{E}^\top \boldsymbol{D}^{-1} \boldsymbol{E}$ through the Woodbury identity, so no dense $m \times m$ inverse is formed. Mean squared error uses the same second-order approximation $g_1 + g_2 + 2g_3 - g_4$ as **sae** (Molina and Marhuenda, 2015), or parametric and non-parametric bootstrap.

Spatio-temporal Fay-Herriot model

When small area data are collected across longitudinal waves $t = 1, \ldots, T$, borrowing strength across space and time yields substantial efficiency gains (Marhuenda et al., 2013): %

$$y_{it} = \boldsymbol{x}_{it}^{\top}\boldsymbol{\beta} + u_{1,it} + u_{2,it} + e_{it}, \tag{6}$$

% where $\boldsymbol{u}_1 = \rho_1 \boldsymbol{W}\boldsymbol{u}_1 + \boldsymbol{\epsilon}_1$ is a spatial SAR effect that is constant over time within each area ($u_{1,it} = u_{1,i}$) and $u_{2,it} = \rho_2 u_{2,i,t-1} + \epsilon_{2,it}$ follows an AR(1) temporal process with $|\rho_2| < 1$, as in Marhuenda et al. (2013). In **fastsae**, the covariance $\boldsymbol{V} = \boldsymbol{A} + \boldsymbol{Z}\boldsymbol{V}_{u_1}\boldsymbol{Z}^{\top}$ (with $\boldsymbol{A}$ block diagonal in $T \times T$ blocks) is never formed: REML scores, Fisher information, and the analytical MSE ($g_1 + g_2 + 2g_3 - g_4$) are computed from $T \times T$, $D \times D$, and low-rank $DT \times (D + p)$ pieces via the Woodbury identity.

Unit-level Battese-Harter-Fuller model

When unit microdata are accessible alongside population auxiliary totals, **fastsae** implements the nested error regression model (Battese et al., 1988): %

$$y_{dj} = \boldsymbol{x}_{dj}^{\top}\boldsymbol{\beta} + u_d + e_{dj}, \quad d = 1, \ldots, D, \quad j = 1, \ldots, n_d, \tag{7}$$

% yielding EBLUP predictors $\hat{\bar{Y}}_d = f_d \bar{y}_d + (1 - f_d)(\bar{\boldsymbol{X}}_{d,r}^{\top}\hat{\boldsymbol{\beta}} + \hat{\gamma}_d(\bar{y}_d - \bar{\boldsymbol{x}}_d^{\top}\hat{\boldsymbol{\beta}}))$, where $f_d = n_d / N_d$.

2.2 Bayesian non-Gaussian hierarchical EBP via INLA

Latent Gaussian models

For non-Gaussian indicators, **fastsae** formulates small area estimation within the Latent Gaussian Model (LGM) framework (Rue et al., 2009): %

$$y_i \mid \eta_i, \boldsymbol{\theta} \sim \pi(y_i \mid \eta_i, \boldsymbol{\theta}), \quad \eta_i = g(\mu_i) = \boldsymbol{x}_i^{\top}\boldsymbol{\beta} + u_i + v_i, \tag{8}$$

% where $g(\cdot)$ is a monotone link function, $v_i \overset{\text{iid}}{\sim} \mathcal{N}(0, \tau_v^{-1})$ is an unstructured effect, and $\boldsymbol{u} \sim \mathcal{N}(\boldsymbol{0}, \boldsymbol{Q}^{-1}(\boldsymbol{\theta}_s))$ is a spatial Gaussian Markov Random Field (GMRF). By default, **fastsae** utilizes the Full Laplace approximation (`strategy = "laplace"`) rather than the simplified Laplace expansion, guaranteeing high analytical fidelity even near boundary domains.

Exact survey dispersion and variance matching

To preserve design-based sampling properties in a Bayesian GLMM, **fastsae** injects distribution-specific dispersion parameters: %

1. **Gaussian**: Precision is fixed to the inverse design variance: $\tau_i = 1/D_i$.
2. **Beta**: For proportions $y_i \in (0, 1)$ with logit link $\text{logit}(\mu_i) = \eta_i$, the Beta variance is $\text{Var}(y_i) = \mu_i(1 - \mu_i)/(1 + \phi_i)$. Equating this to design variance D_i yields the Janicki (2020) dispersion:

$$\phi_i = \max\left(\frac{y_i(1 - y_i)}{D_i} - 1,\, 1\right). \tag{9}$$

3. **Gamma**: For positive skewed expenditures $y_i > 0$ with log link $\log(\mu_i) = \eta_i$, equating $\text{Var}(y_i) = \mu_i^2 / s_i$ to D_i yields the shape precision parameter:

$$s_i = \frac{y_i^2}{D_i} = \frac{1}{\text{CV}_i^2}. \tag{10}$$

4. **Binomial, Poisson, and Negative Binomial**: Domain sample sizes n_i map to `Ntrials`, and population exposures E_i map to the log-link offset $\log(E_i)$.

Spatial priors and space-time interactions

The spatial random field $\boldsymbol{u}$ supports five graph topologies: Besag ICAR (Besag et al., 1991), scaled BYM2 with Penalized Complexity (PC) priors (Riebler et al., 2016; Simpson et al., 2017), Leroux CAR with exact precision $\boldsymbol{Q} = \tau[(1 - \rho)\boldsymbol{I} + \rho(\text{diag}(\boldsymbol{d}) - \boldsymbol{W})]$ (Leroux et al., 2000), and the Spatial Lag

Model (SLM). For panel observations across $T \geq 2$, temporal trends support random walks (RW1, RW2) and AR(1) dynamics. Space-time interactions can be specified as additive, domain-specific, separable, or according to the four Knorr-Held (2000) interaction types ($\boldsymbol{Q}_{st} = \boldsymbol{Q}_s \otimes \boldsymbol{Q}_t$).

2.3 Two-stage hierarchical benchmarking calibration

In official statistics, model-based small area estimates $\hat{\theta}_{di}$ for domain i in province/region g must aggregate to officially published higher-level targets (Rao and Molina, 2015; Stefan and Hidiroglou, 2020). Let w_{di} denote domain population weights such that $W_g = \sum_{i \in g} w_{di}$ is the provincial total weight, and $W_{\text{nat}} = \sum_g W_g$ is the national total weight. Let T_g denote the provincial target and T_{nat} denote the national aggregate target.

fastsae implements a two-stage hierarchical benchmarking calibration: %

1. **Stage 1: Harmonization of Provincial Targets to National Benchmark**: When initial provincial targets T_g (or initial aggregated model estimates $\hat{\tilde{\theta}}_g = \sum_{i \in g}(w_{di}/W_g)\hat{\theta}_{di}$) do not sum to the national target T_{nat}, calibrated provincial targets $\tilde{T}_g$ are computed to satisfy $\sum_g (W_g/W_{\text{nat}})\tilde{T}_g = T_{\text{nat}}$.
2. **Stage 2: Domain-Level Calibration within Provinces**: Domain estimates within province g are calibrated to satisfy $\sum_{i \in g}(w_{di}/W_g)\hat{\theta}^{\text{BM}}_{di} = \tilde{T}_g$.

% Across both stages, **fastsae** supports four calibration mechanisms: %

- **Ratio Benchmarking** (You and Rao, 2002): $\hat{\theta}^{\text{BM}}_{di} = \hat{\theta}_{di}\left(\tilde{T}_g / \sum_{j \in g} \tilde{w}_{dj}\hat{\theta}_{dj}\right)$.
- **Difference Benchmarking** (Rao and Molina, 2015): $\hat{\theta}^{\text{BM}}_{di} = \hat{\theta}_{di} + \left(\tilde{T}_g - \sum_{j \in g} \tilde{w}_{dj}\hat{\theta}_{dj}\right)$.
- **Variance-Weighted Optimal Benchmarking** (Datta et al., 2011; Steorts and Ghosh, 2014): Adjustments minimize $\sum_i (\hat{\theta}^{\text{BM}}_{di} - \hat{\theta}_{di})^2/\text{MSE}(\hat{\theta}_{di})$ subject to the benchmark constraint, which allocates them in proportion to $\tilde{w}_{di}\,\text{MSE}(\hat{\theta}_{di})$:

$$\hat{\theta}^{\text{BM}}_{di} = \hat{\theta}_{di} + \tilde{w}_{di}\text{MSE}(\hat{\theta}_{di}) \left(\sum_{j \in g} \tilde{w}^2_{dj}\text{MSE}(\hat{\theta}_{dj})\right)^{-1} \left(\tilde{T}_g - \sum_{j \in g} \tilde{w}_{dj}\hat{\theta}_{dj}\right). \tag{11}$$

- **Logit-Scale Benchmarking**: For bounded proportions in $(0,1)$, linear adjustments can push estimates outside $[0,1]$. **fastsae** solves for an additive shift α_g on the logit scale:

$$\sum_{i \in g} \tilde{w}_{di}\,\text{logit}^{-1}\left(\text{logit}(\hat{\theta}_{di}) + \alpha_g\right) = \tilde{T}_g, \tag{12}$$

  via robust one-dimensional root-finding, guaranteeing $\hat{\theta}^{\text{BM}}_{di} \in (0,1)$.

2.4 Official statistics reliability classification

NSOs classify domain estimates into publication tiers based on the Relative Standard Error ($\text{RSE} = \frac{\sqrt{\text{MSE}(\hat{\theta}_i)}}{|\hat{\theta}_i|} \times 100\%$). **fastsae** provides `add_reliability_flags()` implementing institutional thresholding: %

$$\text{Tier}_i = \begin{cases} \text{Reliable (Direct Publish)}, & \text{RSE}_i \leq c_1, \\ \text{Use with Caution (Accompanying Warning)}, & c_1 < \text{RSE}_i \leq c_2, \\ \text{Unreliable (Suppress or Aggregate)}, & \text{RSE}_i > c_2, \end{cases} \tag{13}$$

% where the default thresholds $(c_1, c_2) = (25\%, 50\%)$ follow the criteria of BPS-Statistics Indonesia (with $\text{RSE} \leq c_1$ reliable and $\text{RSE} > c_2$ unreliable); other agencies' cut-offs can be supplied.

%% – Software Design ————————————————————-

3 Software design and implementation

The architecture of **fastsae** unites compiled linear algebra with high-level Bayesian GMRF modeling and production reporting. Figure 1 outlines the structural flow.

The frequentist engines (`eblup_fh()`, `eblup_sfh()`, `eblup_stfh()`, `eblup_bhf()`, `eblup_twofold()`) never form a dense inverse of the marginal covariance: the spatial model works with sparse Cholesky

fastsae High-Level Software Architecture

1. User-Facing Production API (R)		
Frequentist Modeling	**Bayesian HB Engine**	**Production Workflow**
`eblup_fh()`	`hb_area(...,family = "beta")`	`benchmark_sae()` (Two-Stage)
`eblup_sfh()`	`hb_area(...,family = "gamma")`	`compare_sae()`
`eblup_stfh()`	`hb_area(...,family = "poisson")`	`add_reliability_flags()`
`eblup_bhf()`	`hb_area(...,spatial = "bym2")`	`map_sae()` / `export_sae()`
2. Core Computational Backends		
Compiled C++17 Core	**INLA Deterministic Engine**	**Spatial Linear Algebra**
RcppArmadillo Matrix Algebra	Sparse Cholesky Factorization	Row-normalized Proximity
OpenMP Threaded Bootstrap	Full Laplace Approximation	PC Priors & GMRF Graphs
In-Place Fisher-Scoring Updates	Survey Dispersion Injections	Spectral Eigenvalue Bounds

Figure 1: Structural architecture and functional components of fastsae.

factors (RcppEigen), the spatio-temporal and two-fold models with block-structured Woodbury identities, and the unit-level model with domain-level sufficient statistics, profiling $\boldsymbol{\beta}$ and σ_e^2 out of the (restricted) likelihood. The bootstrap MSE loops run in parallel with OpenMP (`#pragma omp parallel for`); random numbers are drawn sequentially from R's generator beforehand, so results depend only on the seed and not on the number of threads. The Bayesian functions (`hb_area()`, `hb_unit()`, `hb_twofold()`) translate formulas, spatial adjacency matrices $\boldsymbol{W}$, and survey dispersion arguments into **R-INLA** calls. How a model is estimated is specified separately from the model by an engine object: `inla_engine()` holds the INLA strategy, priors, and further `INLA::inla()` arguments, and `jax_engine()` fits `hb_area()` models by NUTS sampling in JAX/NumPyro (CPU or GPU) through reticulate. Standard S3 methods (`print()`, `summary()`, `plot()`, `coef()`, `fitted()`, `residuals()`) are implemented across all model classes.

%% – Computational Benchmarks ————————————————-

4 Computational benchmarks and simulations

4.1 Frequentist scalability and memory efficiency: fastsae vs. sae and emdi

We evaluated computational speed and memory allocation across domain sizes $D = 30, 50, 100, 250, 500,$ and 1000 using simulated datasets with 5 auxiliary covariates. Times are medians measured with the bench package; memory is the total memory allocated by R during one call (`mem_alloc`), not the peak resident memory. The fastsae rows were measured with fastsae 1.0.0 under R 4.6.1 on an Apple M4 processor (10 cores); the sae and emdi rows come from an earlier run of the same design on different random draws. The scripts are in the `benchmarks/` directory of the package repository.

Table 2 demonstrates the scalability gains of fastsae. At $D = 1,000$, classical FH completes in about one millisecond in fastsae compared to 1.5 seconds in sae ($1,263\times$ slower) and 51 seconds in emdi ($42,874\times$ slower). For spatio-temporal models at $D = 1,000$ ($T = 5$), sae required nearly 30 minutes (1,786 seconds) and allocated 32.5 gigabytes, whereas fastsae finished in 75 seconds allocating 3.3 megabytes. **Operational Takeaway**: fastsae eliminates out-of-memory bottlenecks, enabling routine spatio-temporal estimation across thousands of domains on standard laptop hardware.

4.2 Bayesian Beta Hierarchical Modeling scalability: fastsae vs. tipsae

For bounded rates, we benchmarked `hb_area(...,family = "beta")` against tipsae (De Nicolò and Gardini, 2022) using Stan (4 chains, 2,000 iterations).

As shown in Table 3, deterministic Laplace approximations in fastsae scale flatly with domain count: fitting a spatial Beta model at $D = 1,000$ requires only 1.87 seconds in fastsae compared to 404 seconds in tipsae ($217\times$ speedup). Most of the fastsae time is the fixed start-up cost of the INLA program (about 1.2 seconds even at $D = 30$). In spatio-temporal panel models ($D = 50, T = 5$), tipsae required 15.1 minutes (908 seconds), whereas fastsae completed in 1.71 seconds ($530\times$ speedup). Despite this speed differential, inferential concordance on $D = 1,000$ yields Pearson correlation $r = 0.9993$, Spearman $\rho = 0.9992$, and Mean Absolute Difference MAD $= 0.0197$. **Operational Takeaway**: Deterministic Laplace approximation in fastsae achieves inferential equivalence with fully

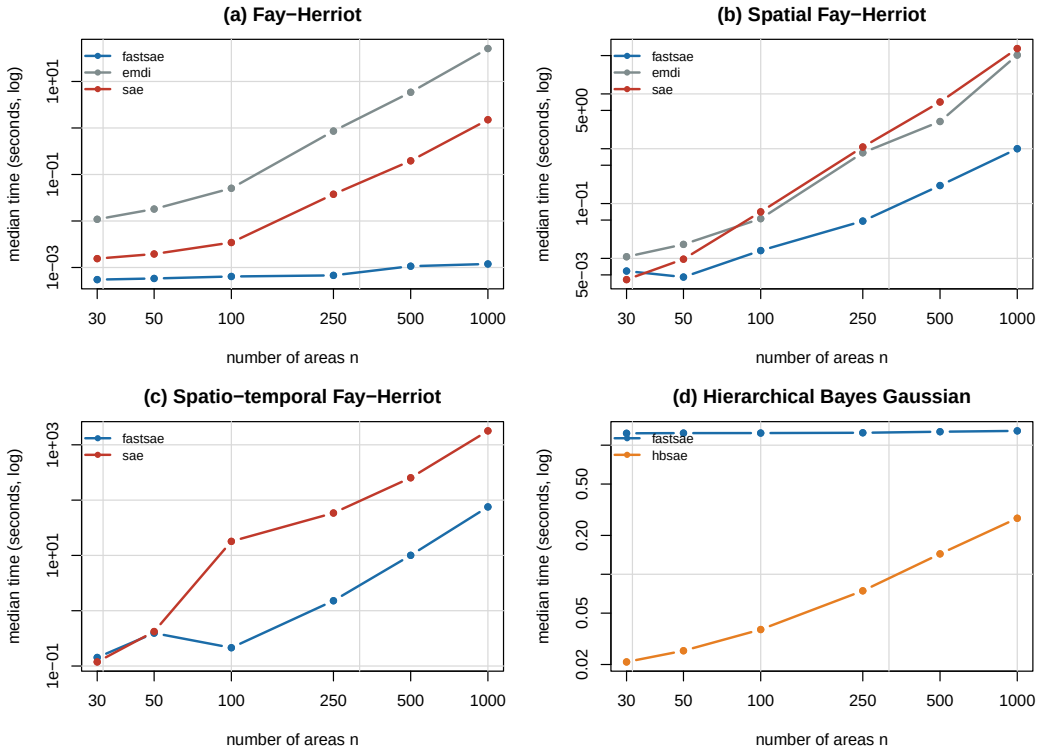


Figure 2: Median runtime scaling across domain sample sizes (D) on log-log scale for classical, spatial, spatio-temporal Fay-Herriot, and Bayesian Gaussian models.

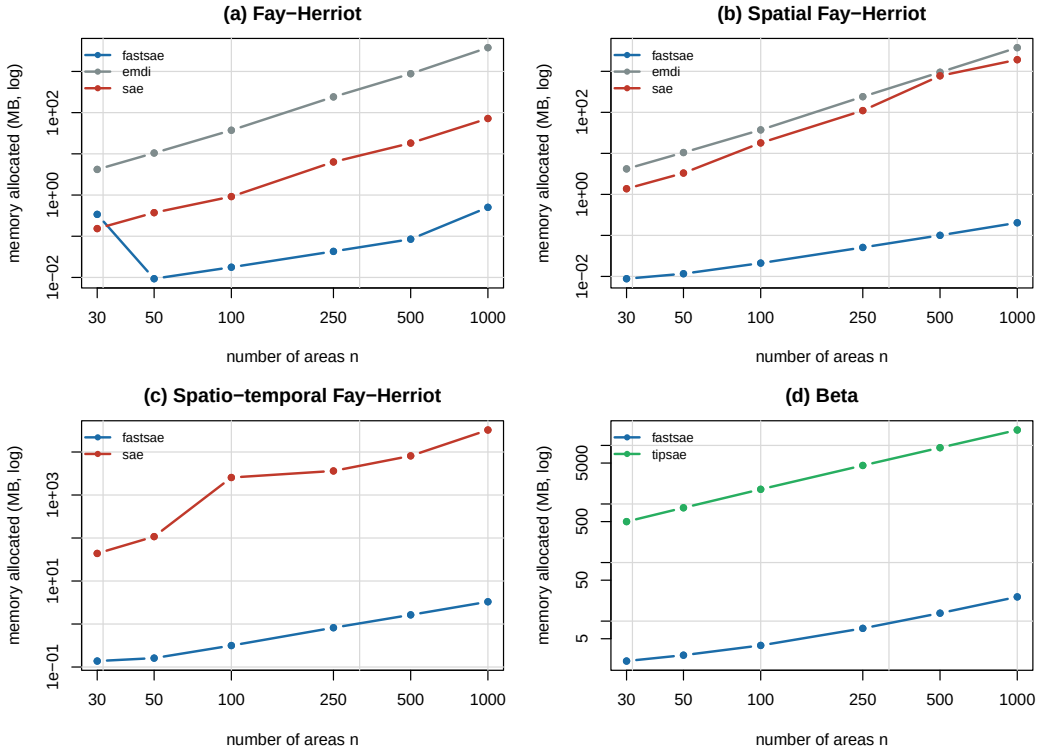


Figure 3: Memory allocated per call (MB, `bench::mark()$mem_alloc`) across domain counts (D) on log-log scale. **fastsae** avoids quadratic $O(D^2T^2)$ allocation.

D	Execution Time: Median (seconds)			Memory Allocated (MB)		
	fastsae	sae	emdi	fastsae	sae	emdi
Fay-Herriot EBLUP (Classical)						
30	0.00055	0.00155	0.01088	0.339	0.153	4.18
100	0.00064	0.00344	0.05051	0.018	0.918	37.34
500	0.00107	0.19673	5.86366	0.085	18.239	880.31
1000	0.00119	1.50363	51.0224	0.503	71.977	3,773.60
Spatial Fay-Herriot (SEBLUP)						
30	0.00588	0.00410	0.01071	0.009	1.373	4.18
100	0.01388	0.07054	0.05299	0.021	17.976	37.34
500	0.21310	7.11793	3.13662	0.101	776.407	949.86
1000	0.99977	67.1518	51.0593	0.203	1,920.01	3,772.57
Spatio-Temporal Fay-Herriot (STEBLUP, $T = 5$)						
30	0.14272	0.11856	–	0.138	43.801	–
100	0.21415	17.9425	–	0.316	2,551.10	–
500	10.0430	253.343	–	1.633	8,124.37	–
1000	75.3340	1,786.49	–	3.293	32,478.3	–

Table 2: Frequentist runtime and memory allocation (point estimates and analytical MSE) across domain sample sizes (D). At $D = 1,000$, fastsae is $1,263\times$ faster than sae for FH, $67\times$ for SFH, and $24\times$ for ST-FH. At $D = 30$ the spatial and spatio-temporal fits are dominated by fixed costs, and sae is slightly faster.

D	Cross-Sectional Beta		Spatial Beta (Besag)		Spatio-Temporal Beta	
	fastsae	tipsae	fastsae	tipsae	fastsae	tipsae
30	1.20 s	2.35 s	1.22 s	25.72 s	1.48 s	540.57 s
50	1.23 s	3.77 s	1.23 s	65.72 s	1.71 s	908.17 s
100	1.26 s	8.24 s	1.25 s	100.02 s	–	–
250	1.32 s	26.28 s	1.33 s	155.90 s	–	–
500	1.40 s	64.21 s	1.46 s	248.82 s	–	–
1000	1.61 s	187.16 s	1.87 s	404.34 s	–	–

Table 3: Median execution time for Bayesian Beta models (spatio-temporal: $T = 5$). Both packages are fitted to the same simulated data. Point estimates of fastsae and tipsae for the cross-sectional model at $D = 1,000$ achieve $r = 0.9993$ concordance.

Estimator	Emp. RMSE	RRMSE (%)	MARB (%)	95% CrI Cov.	CrI Width	Runtime
Direct Sample Estimator	0.0719	27.1%	21.7%	–	–	–
fastsae (Spatial Besag)	**0.0481**	**19.1%**	16.9%	89.0%	**0.173**	**4.17 s**
tipsae (Spatial Besag)	0.0487	18.7%	**15.7%**	**93.7%**	0.202	36.86 s

Table 4: Finite population Monte Carlo simulation results ($D = 500$, $R = 25$). fastsae achieves an efficiency gain of $1.92\times$ over the direct survey estimator.

converged Hamiltonian MCMC in seconds rather than minutes or hours.

4.3 Monte Carlo finite population validation against ground truth

To assess frequentist design properties, we conducted design-based Monte Carlo simulations under known ground truth at national scale ($D = 500$ small areas, $R = 25$ independent sample draws) using a synthetic finite population of $N = 250,000$ individuals.

Table 4 demonstrates that `hb_area()` reduces empirical RMSE from 0.0719 (direct survey estimator) to 0.0481, yielding a $1.92\times$ empirical efficiency gain. Relative Root Mean Squared Error drops from 27.1% to 19.1%. Comparing Bayesian engines, fastsae achieves a slightly lower empirical RMSE than tipsae (0.0481 vs. 0.0487) and produces narrower credible intervals (0.173 vs. 0.202) while requiring nearly $9\times$ less runtime per fit. **Operational Takeaway**: fastsae produces well-calibrated, highly efficient area-level predictions that closely track true finite-population parameters while scaling effortlessly to national-scale administrative systems.

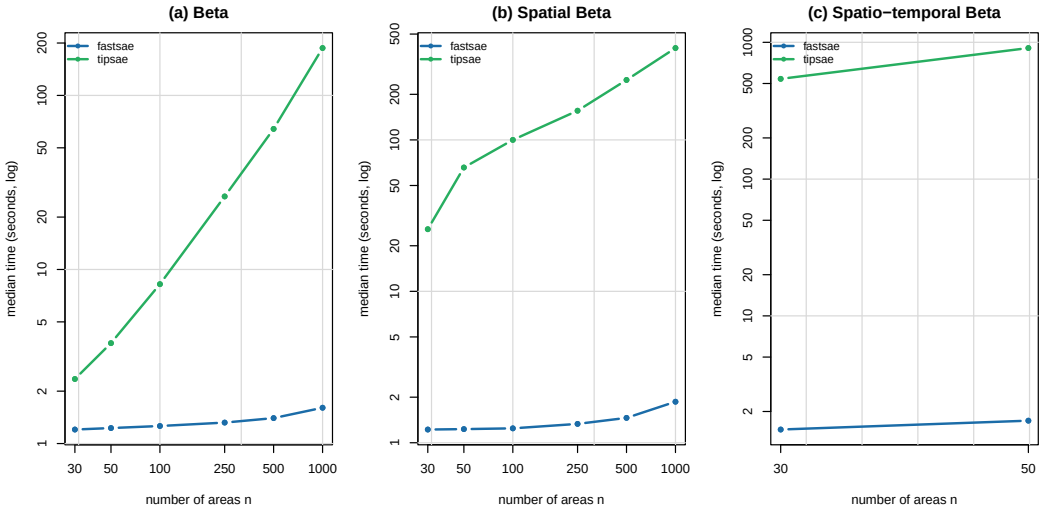


Figure 4: Runtime scaling across domain sample sizes (D) for cross-sectional, spatial, and spatio-temporal Bayesian Beta models.

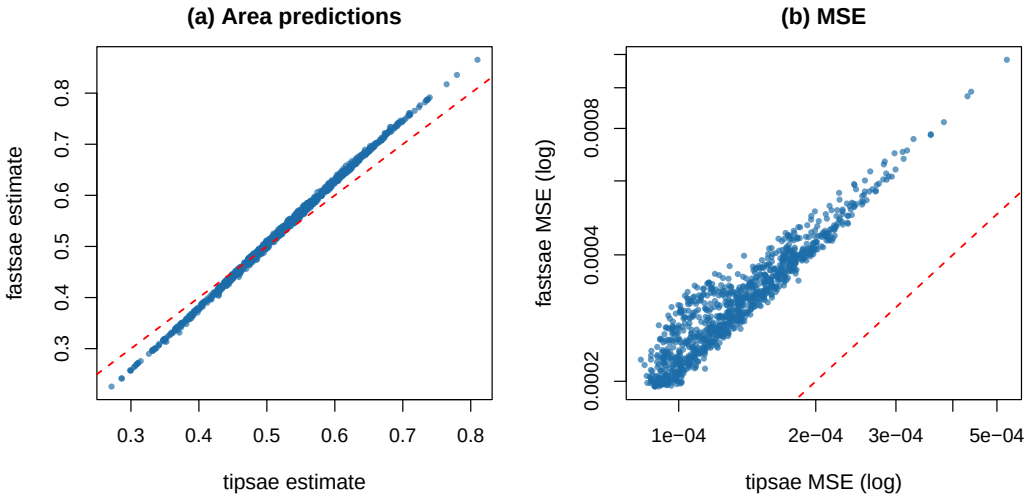


Figure 5: Numerical equivalence and inferential concordance between **fastsae** (INLA) and **tipsae** (Stan HMC) at $D = 1{,}000$ ($r = 0.9993$).

%% – Applied Illustration ——————————————————

5 Applied illustration and production workflow

We demonstrate an end-to-end NSO production workflow using the benchmark dataset `sim_area` ($D = 42$ domains across 6 provincial groups).

5.1 Loading data, spatial weights, and fitting models

```
library("fastsae")
data("sim_area")
sim_area$prov <- rep(paste0("Prov_", 1:6), each = 7)
sim_area$pop <- round(runif(nrow(sim_area), 15000, 45000))

# Construct row-standardized spatial proximity matrix
coords <- cbind(sim_area$x_coord, sim_area$y_coord)
dists <- as.matrix(dist(coords))
W <- (dists < quantile(dists, 0.18)) * 1
diag(W) <- 0
rs <- rowSums(W); rs[rs == 0] <- 1; W_std <- W / rs

# Fit Frequentist and Bayesian Models
fit_fh <- eblup_fh(y_gaussian ~ x1 + x2, vardir = "vardir",
  data = sim_area, method = "REML", print_result = FALSE)
fit_sfh <- eblup_sfh(y_gaussian ~ x1 + x2, vardir = "vardir",
  domain = "area", W = W_std, data = sim_area,
  method = "REML", print_result = FALSE)
fit_beta <- hb_area(y_beta ~ x1 + x2, vardir = "vardir",
  data = sim_area, family = "beta", spatial = "bym2",
  W = W, print_result = FALSE)
```

5.2 Automated model concordance evaluation

Practitioners can immediately compare competitive models via `compare_sae()`:

```
comp <- compare_sae(fit_fh, fit_sfh, names = c("FH", "SFH"))
comp$metrics[c(1, 2, 4, 6, 8, 9), ]
```

```
#>                                                Metric  Value
#> 1                                     Matched Domains     42
#> 2                              Pearson Correlation (r) 0.9996
#> 4                       Mean Absolute Difference (MAE) 0.0279
#> 6                            Mean MSE Ratio (FH / SFH) 0.9897
#> 8                      Domains where SFH has lower MSE 3 / 42
#> 9 Percentage of domains where SFH is more efficient (%)   7.1%
```

5.3 Two-stage hierarchical benchmarking calibration

To harmonize domain estimates with an official national poverty target ($T_{\text{nat}} = 0.50$), we apply two-stage calibration:

```
bm <- benchmark_sae(fit_fh, national_target = 0.50, group = "prov",
  weight = "pop", method = "ratio", type = "mean")
head(bm[, c("domain", "group", "original", "benchmarked", "rel_adjustment_pct")], 4)
```

```
#> === fastsae Small Area Benchmark Calibration ===
#> Method: Ratio (Multiplicative Adjustment)
#> Target Type: Weighted Mean / Rate
#> Total Domains: 4
#>
```

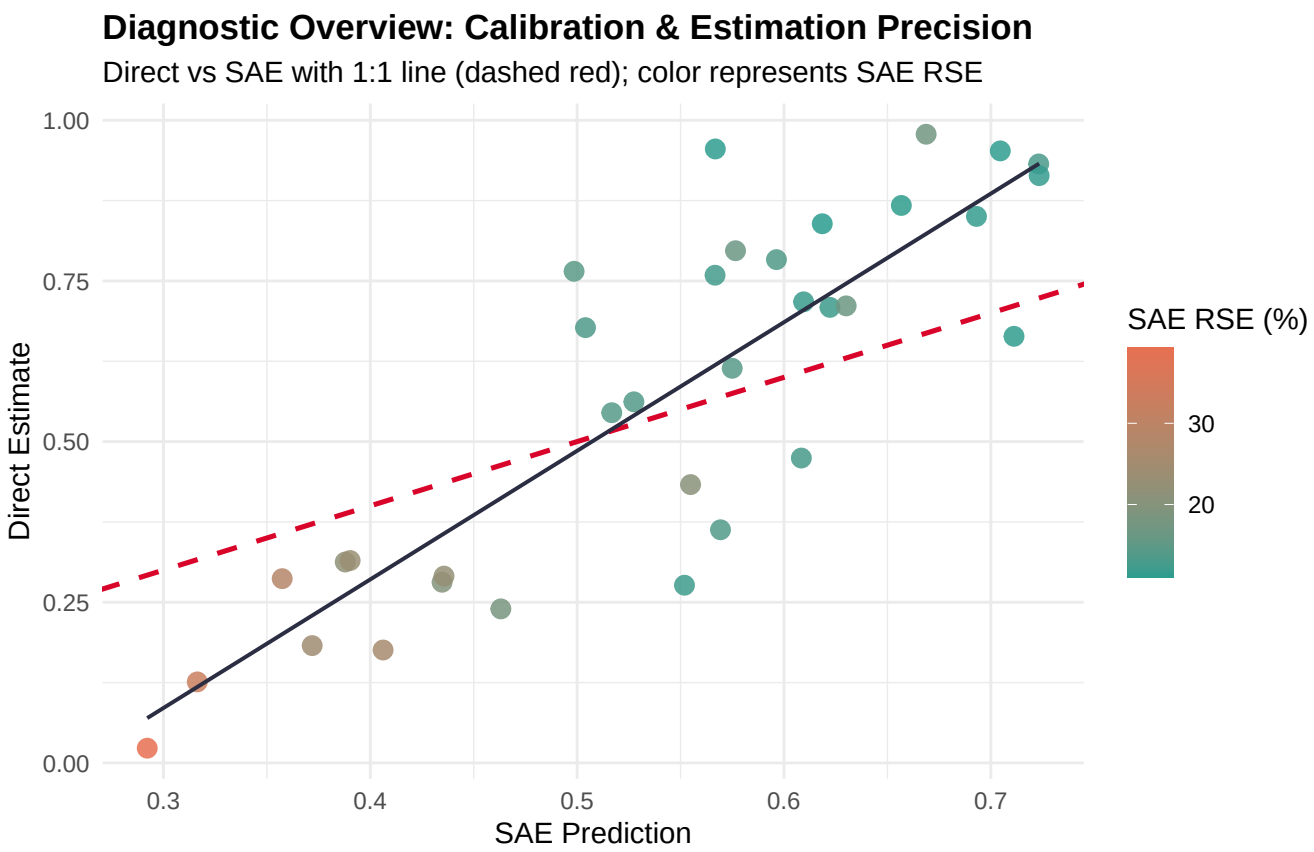


Figure 6: Model diagnostic panels generated via autoplot.

```
#> First 6 benchmarked domains (Level 2):
#>   domain  group   original benchmarked rel_adjustment_pct
#> 1      1 Prov_1  2.6646558   0.8916957          -66.53618
#> 2      2 Prov_1 -1.6547233  -0.5537337          -66.53618
#> 3      3 Prov_1  1.2504674   0.4184542          -66.53618
#> 4      4 Prov_1 -0.7988192  -0.2673155          -66.53618
```

The calibrated estimates satisfy $\sum(w_i / \sum w_i)\hat{\theta}_i^{\mathrm{BM}} = 0.50$ exactly.

5.4 Official statistics reliability classification

Using `add_reliability_flags()`, domains are classified into official dissemination tiers:

```
flags <- add_reliability_flags(fit_beta)
table(flags$reliability)
```

```
#>
#>          Reliable (< 20%) Use with Caution (20-30%)        Unreliable (≥ 30%)
#>                        31                          8                         3
```

5.5 Model diagnostics and visualization

Diagnostics via `diagnose()` and `autoplot()` provide visual quality verification:

```
diag_beta <- diagnose(fit_beta)
autoplot(diag_beta)
```

%% – Summary and Discussion ————————————————

Dimension	Observed Empirical Evidence	Institutional Implication ("So What")
Frequentist Scalability	$1,263\times$ speedup for FH and $67\times$ for SFH at $D = 1,000$; a few MB allocated for ST-FH (vs 32 GB in **sae**).	Eliminates server memory exhaustion, allowing desktop computers to execute nationwide bootstrap MSE routines.
Bayesian EBP via INLA	$117\times$ to $530\times$ speedup over Stan MCMC at the largest benchmarked sizes; point correlation $r = 0.9993$ on $D = 1,000$.	Eliminates manual MCMC chain diagnostics and divergent transition failures in routine production timelines.
Design-Based Ground Truth	$1.92\times$ empirical efficiency gain over direct estimator; 89% empirical coverage.	Delivers substantial MSE reductions while maintaining nominal interval coverage under repeated sampling.
Hierarchical Benchmarking	Simultaneous Level 0 and Level 1 target alignment with logit boundary safety.	Guarantees numeric consistency with published aggregates, eliminating public confusion without violating probability bounds.
Turnkey Dissemination	Direct RSE classification, spatial choropleths, and multi-sheet exports.	Collapses multi-script analyst workflows into standardized, reproducible publication pipelines.

Table 5: Synthesis of empirical findings and operational implications of **fastsae**.

6 Summary and discussion

Small area estimation in official statistics requires reconciling mathematical rigor with strict computational and operational constraints. **fastsae** resolves these challenges by bridging high-performance compiled linear algebra, deterministic Bayesian approximations, and production-ready dissemination tools.

As synthesized in Table 5, **fastsae** provides substantial efficiency gains across both frequentist and Bayesian paradigms. Future developments will focus on expanding unit-level spatial EBP under non-Gaussian likelihoods, integrating zero-inflated Beta (ZIB) specifications for extreme sparse indicators, and supporting direct spatial polygon streaming.

%% – Computational Details & Acknowledgments ————————————

Computational details

The **fastsae** results in this paper were obtained using R 4.6.1 with **fastsae** 1.0.0, **R-INLA** 26.8.7, and **ggplot2** 4.0.3 on an Apple M4 processor; the comparison runs used **sae** 1.3, **emdi** 2.2, **tipsae** 1.0, and **hbsae** 1.2. R and packages from CRAN are available at `https://CRAN.R-project.org/`. **R-INLA** is available from `https://www.r-inla.org/`. The **fastsae** package is available on GitHub at `https://github.com/ridsonap/fastsae`.

Acknowledgments

The authors thank the Department of Computational Statistics and Department of Statistics at Politeknik Statistika STIS for institutional support. We also acknowledge the valuable domain feedback provided by colleagues at BPS-Statistics Indonesia.

%% – Bibliography ————————————————————————

References

References

G. E. Battese, R. M. Harter, and W. A. Fuller. An error-components model for prediction of county crop areas using survey and satellite data. *Journal of the American Statistical Association*, 83(401):28–36,

1988. doi: 10.1080/01621459.1988.10478561. [p1, 4]

J. Besag, J. York, and A. Mollié. Bayesian image restoration, with two applications in spatial statistics. *Annals of the Institute of Statistical Mathematics*, 43(1):1–20, 1991. doi: 10.1007/BF00116466. [p4]

H. J. Boonstra. **mcmcsae***: Markov Chain Monte Carlo Small Area Estimation*, 2022. URL https://CRAN.R-project.org/package=mcmcsae. R package version 0.7.3. [p2]

M. Boubeta, M. J. Lombardía, and D. Morales. Empirical best prediction under area-level Poisson mixed models. *TEST*, 25(3):540–569, 2016. doi: 10.1007/s11749-015-0468-x. [p2]

J. P. Burgard, J. Krause, and R. Merkle. Computational aspects of small area estimation. *Statistical Papers*, 61(6):2533–2555, 2020. doi: 10.1007/s00362-020-01185-w. [p1]

G. S. Datta, M. Ghosh, R. Steorts, and J. Maples. Bayesian benchmarking with applications to small area estimation. *TEST*, 20(3):574–588, 2011. doi: 10.1007/s11749-010-0218-y. [p5]

S. De Nicolò and A. Gardini. Mapping proportions and indicators on the unit interval in small area estimation with the R package **tipsae**. *The R Journal*, 14(4):166–184, 2022. doi: 10.32614/RJ-2022-049. [p2, 6]

E. Fabrizi, M. R. Ferrante, and C. Trivisano. Hierarchical Bayes approaches to model-based small area estimation with application to the estimation of poverty indicators. *Computational Statistics & Data Analysis*, 55(4):1636–1647, 2011. doi: 10.1016/j.csda.2010.10.015. [p2]

R. E. Fay and R. A. Herriot. Estimates of income for small places: An application of James-Stein procedures to census data. *Journal of the American Statistical Association*, 74(366a):269–277, 1979. doi: 10.1080/01621459.1979.10482505. [p1, 3]

W. González-Manteiga, M. J. Lombardía, I. Molina, D. Morales, and L. Santamaría. Bootstrap mean squared error of a small-area EBLUP. *Journal of Statistical Computation and Simulation*, 78(5):443–462, 2008. doi: 10.1080/00949650601141811. [p2]

M. W. Janicki. Properties and applications of beta regression models for small area estimation. Technical Report Research Report Series (Statistics #2020-02), Center for Statistical Research and Methodology, U.S. Census Bureau, Washington, D.C., 2020. URL https://www.census.gov/content/dam/Census/library/working-papers/2020/adrm/rrs2020-02.pdf. [p2, 4]

L. Knorr-Held. Bayesian modelling of spatio-temporal disease mapping. *Statistics in Medicine*, 19(17–18):2555–2567, 2000. doi: 10.1002/1097-0258(20000915/30)19:17/18<2555::AID-SIM587>3.0.CO;2-#. [p5]

A.-K. Kreutzmann, S. Pannier, N. Rojas-Perilla, T. Schmid, M. Templ, and N. Tzavidis. The R package **emdi** for estimating and mapping indicators to monitor poverty and inequality. *Journal of Statistical Software*, 91(7):1–33, 2019. doi: 10.18637/jss.v091.i07. [p2]

B. G. Leroux, X. Lei, and N. Breslow. Estimation of disease rates in small areas: A new mixed model for spatial dependence. In *Statistical Models in Epidemiology, the Environment, and Clinical Trials*, pages 135–178. Springer-Verlag, New York, 2000. doi: 10.1007/978-1-4612-1284-3_4. [p4]

Z. R. Li, Y. Hsiao, J. Godwin, B. D. Martin, J. Wakefield, and S. J. Clark. Changes in the spatial distribution of the under-five mortality rate: Small-area analysis of 122 DHS surveys in 262 subregions of 35 countries in Africa. *PLOS Medicine*, 17(1):e1003061, 2020. doi: 10.1371/journal.pmed.1003061. [p2]

Y. Marhuenda, I. Molina, and D. Morales. Small area estimation with spatio-temporal Fay-Herriot models. *Computational Statistics & Data Analysis*, 58:308–325, 2013. doi: 10.1016/j.csda.2012.09.002. [p1, 4]

I. Molina and Y. Marhuenda. **sae**: An R package for small area estimation. *The R Journal*, 7(1):73–87, 2015. doi: 10.32614/RJ-2015-007. [p2, 3]

A. Petrucci and N. Salvati. Small area estimation for spatial data in presence of correlated random area effects. *Journal of Official Statistics*, 22(3):407–428, 2006. URL https://www.scb.se/contentassets/ca21efb41fee47d293bbee5bf7be7fb3/small-area-estimation-for-spatial-data-in-presence-of-correlated-random-area-effects.pdf. [p1, 3]

D. Pfeffermann. New important developments in small area estimation. *Statistical Science*, 28(1):40–68, 2013. doi: 10.1214/12-STS395. [p1]

N. G. N. Prasad and J. N. K. Rao. The estimation of the mean squared error of small-area estimators. *Journal of the American Statistical Association*, 85(409):163–171, 1990. doi: 10.1080/01621459.1990.10475320. [p3]

M. Pratesi, editor. *Analysis of Poverty Data by Small Area Estimation*. John Wiley & Sons, Chichester, United Kingdom, 2016. doi: 10.1002/9781118814963. [p1]

M. Pratesi and N. Salvati. Small area estimation: The EBLUP estimator based on spatially correlated random area effects. *Statistical Methods & Applications*, 17(1):113–141, 2008. doi: 10.1007/s10260-007-0061-9. [p1, 3]

J. N. K. Rao and I. Molina. *Small Area Estimation*. John Wiley & Sons, Hoboken, New Jersey, 2nd edition, 2015. doi: 10.1002/9781118735855. [p1, 5]

A. Riebler, S. H. Sørbye, D. Simpson, and H. Rue. An intuitive Bayesian spatial model for disease mapping that accounts for scaling. *Statistical Methods in Medical Research*, 25(4):1145–1165, 2016. doi: 10.1177/0962280216660421. [p4]

H. Rue, S. Martino, and N. Chopin. Approximate Bayesian inference for latent Gaussian models by using integrated nested Laplace approximations. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 71(2):319–392, 2009. doi: 10.1111/j.1467-9868.2008.00700.x. [p2, 4]

D. Simpson, H. Rue, A. Riebler, T. G. Martins, and S. H. Sørbye. Penalising model component complexity: A principled, practical approach to constructing priors. *Statistical Science*, 32(1):1–28, 2017. doi: 10.1214/16-STS576. [p2, 4]

M. Stefan and M. A. Hidiroglou. Benchmarking small area estimators with a two-stage calibration approach. *Survey Methodology*, 46(2):241–262, 2020. URL https://www150.statcan.gc.ca/n1/pub/12-001-x/2020002/article/00007-eng.htm. [p2, 5]

R. C. Steorts and M. Ghosh. On benchmarking with applications to the estimation of mean squared errors of small area estimators. *Biometrika*, 101(2):397–412, 2014. doi: 10.1093/biomet/ast053. [p5]

J. Wang, W. A. Fuller, and Y. Qu. Small area estimation under a restriction. *Survey Methodology*, 34 (1):29–36, 2008. URL https://www150.statcan.gc.ca/n1/pub/12-001-x/2008001/article/10619-eng.pdf. [p2]

Y. You and J. N. K. Rao. Small area estimation using unmatched sampling and linking models. *The Canadian Journal of Statistics*, 30(1):3–15, 2002. doi: 10.2307/3315862. [p2, 5]

Ridson Al Farizal P.
Department of Computational Statistics, Politeknik Statistika STIS
Jl. Otto Iskandardinata No. 64C
Jakarta Timur 13330, Indonesia
ORCiD: 0000-0003-0617-0214
alfrzlp@gmail.com

Azka Ubaidillah
Department of Statistics, Politeknik Statistika STIS
Jl. Otto Iskandardinata No. 64C
Jakarta Timur 13330, Indonesia
ORCiD: 0000-0002-3597-0459
azka@stis.ac.id